

Retinal Vessel Segmentation

**Image Segmentation: Exploring
U-Nets**

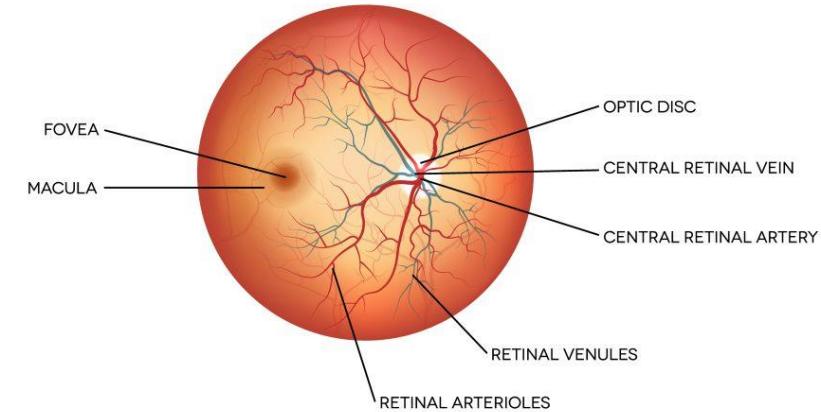
Content

- DRIVE dataset
- UNet applications
- Architectures and findings
- Final results and visualizations

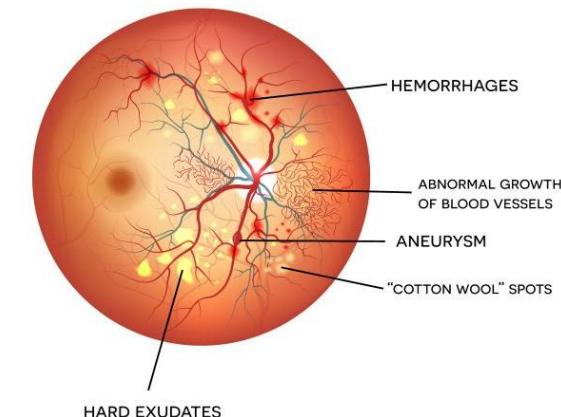
Why?

- Early detection of diseases
 - diabetic retinopathy, hypertension, ...
- Assess overall vascular health
- Non-invasive
 - Reduced risk of complications compared to surgical methods
- Cost-effective
 - Reduced cost compared to surgical methods

NORMAL RETINA

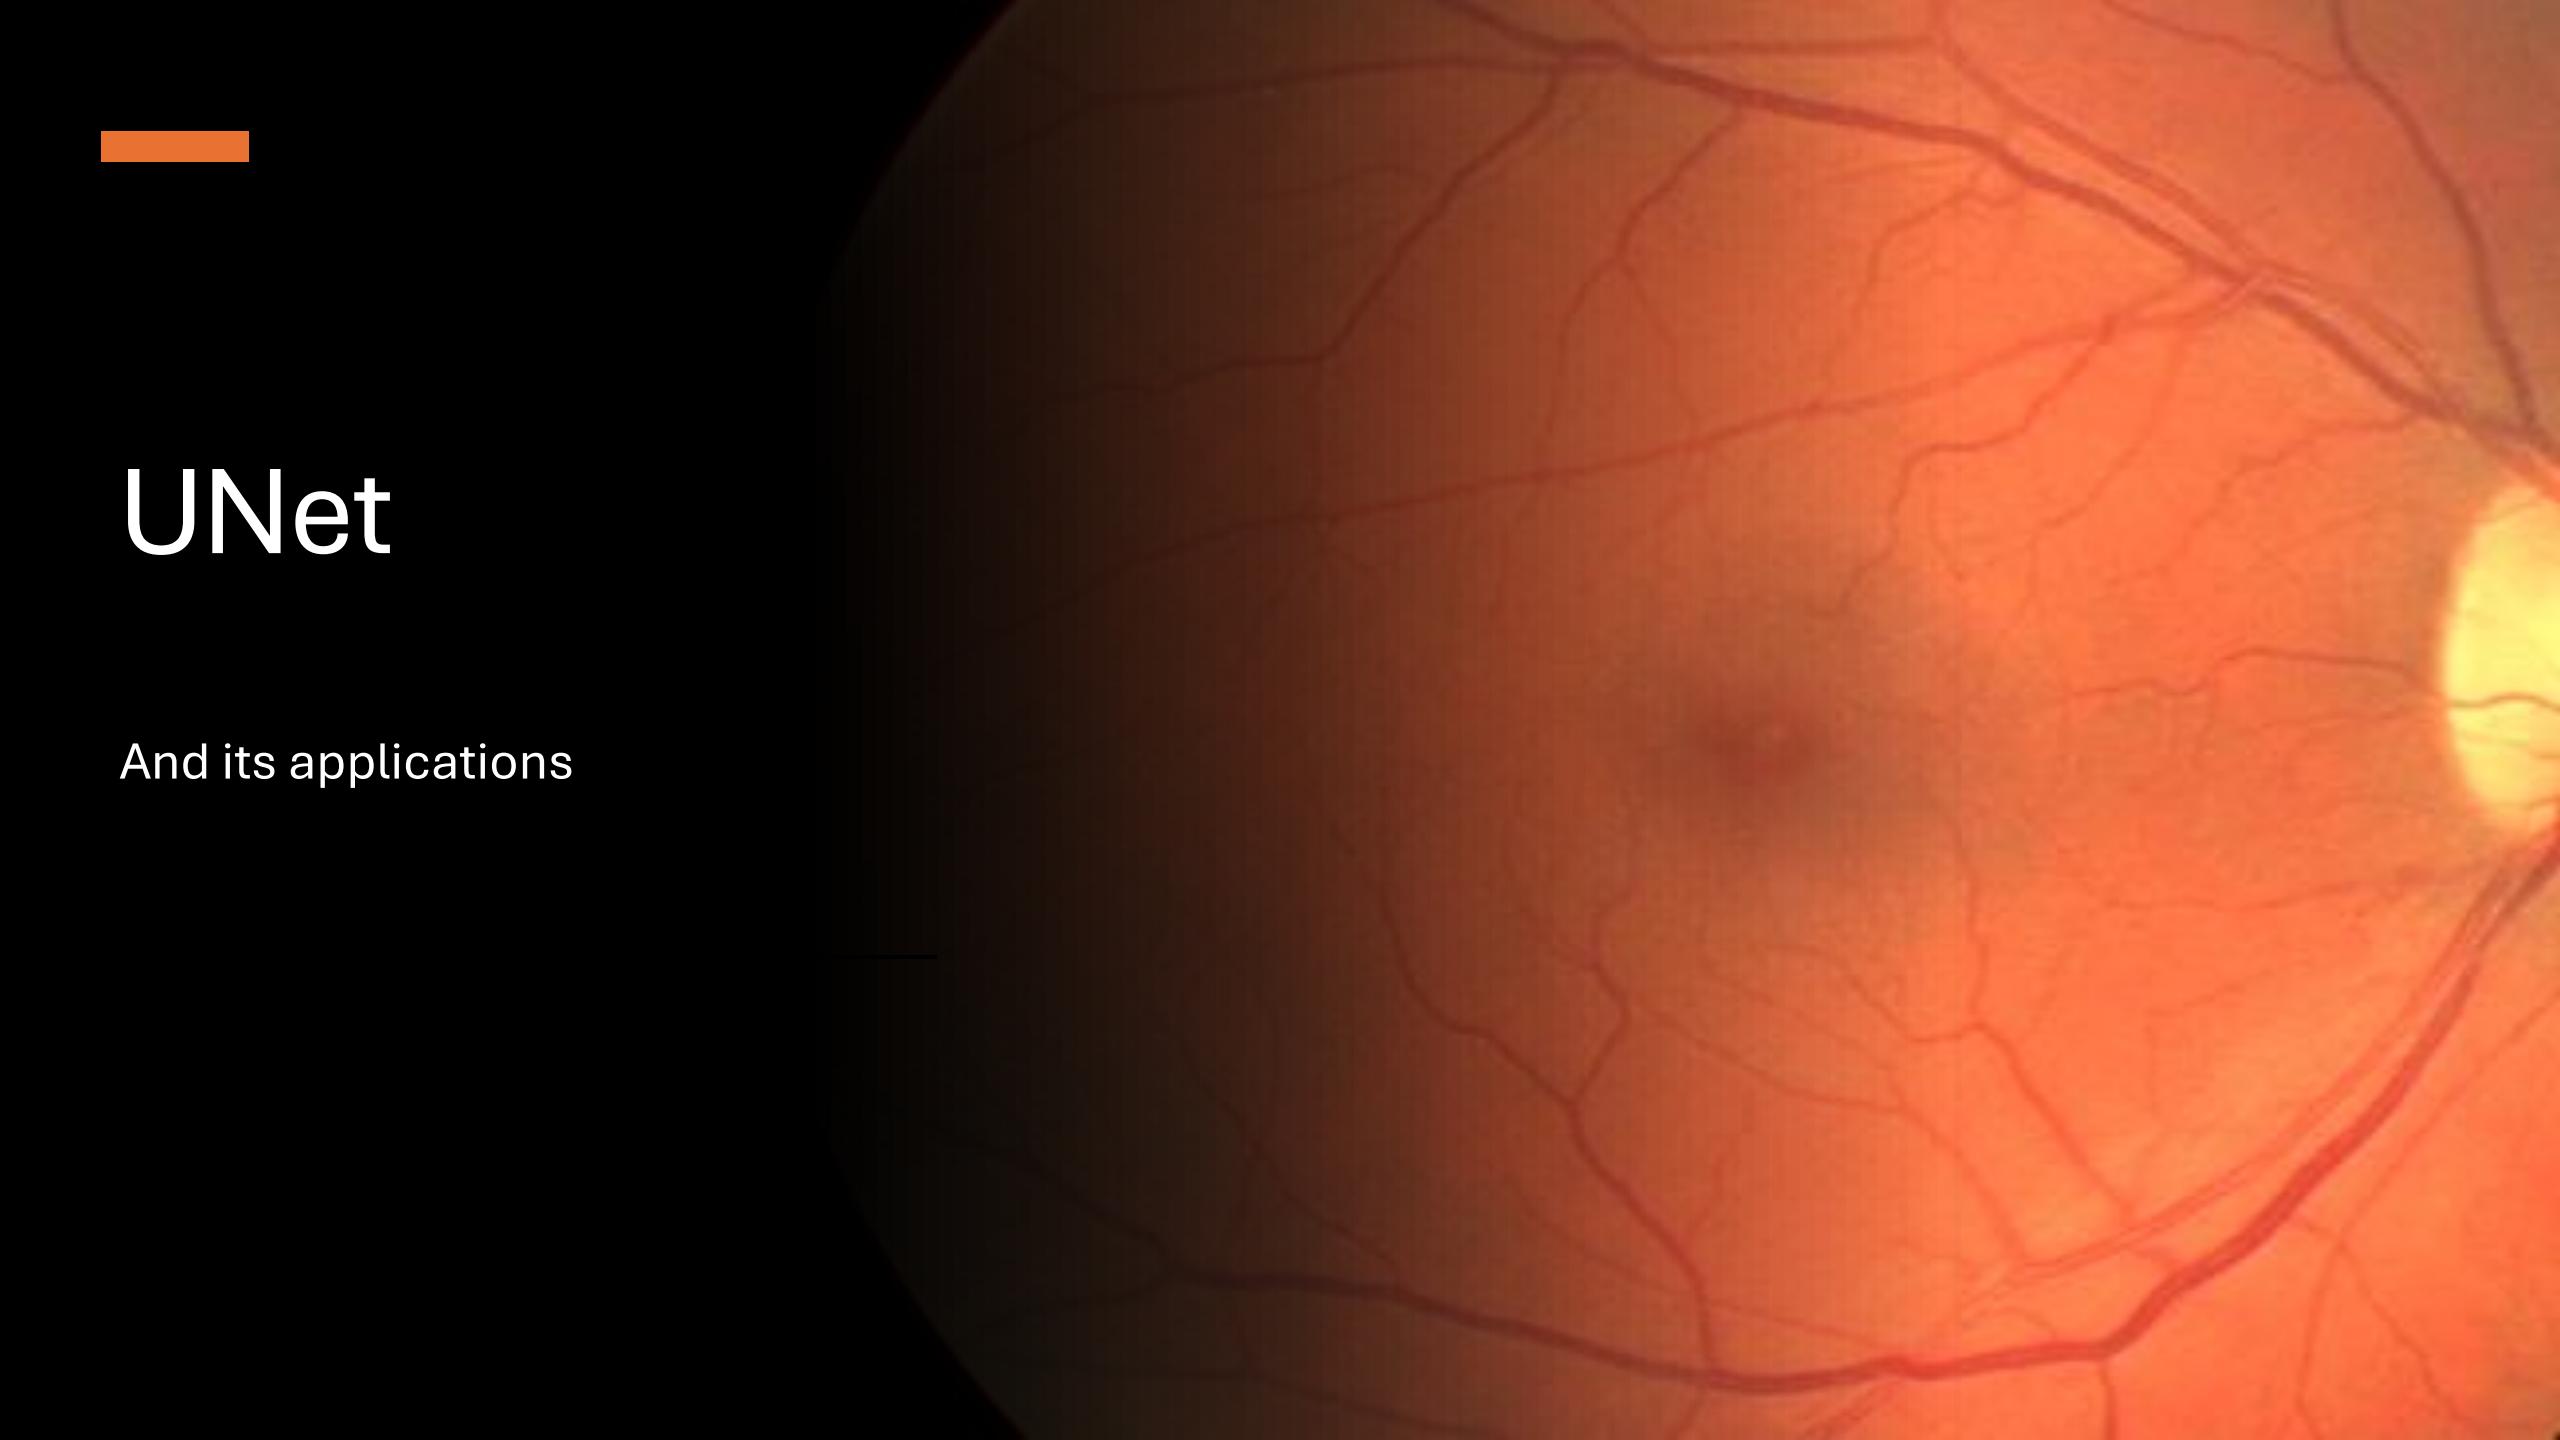


DIABETIC RETINOPATHY



-
- Goal:
 - Segmentation of retinal blood vessels
 - Challenges:
 - Only 20 labeled images
 - Highly imbalanced data
 - Thick vessels <-> fine vessels
 - Noise
-

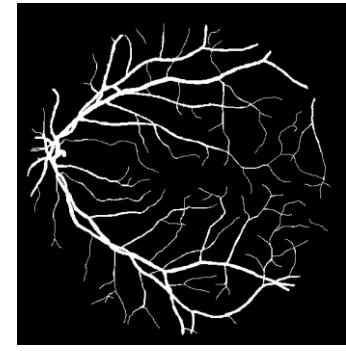
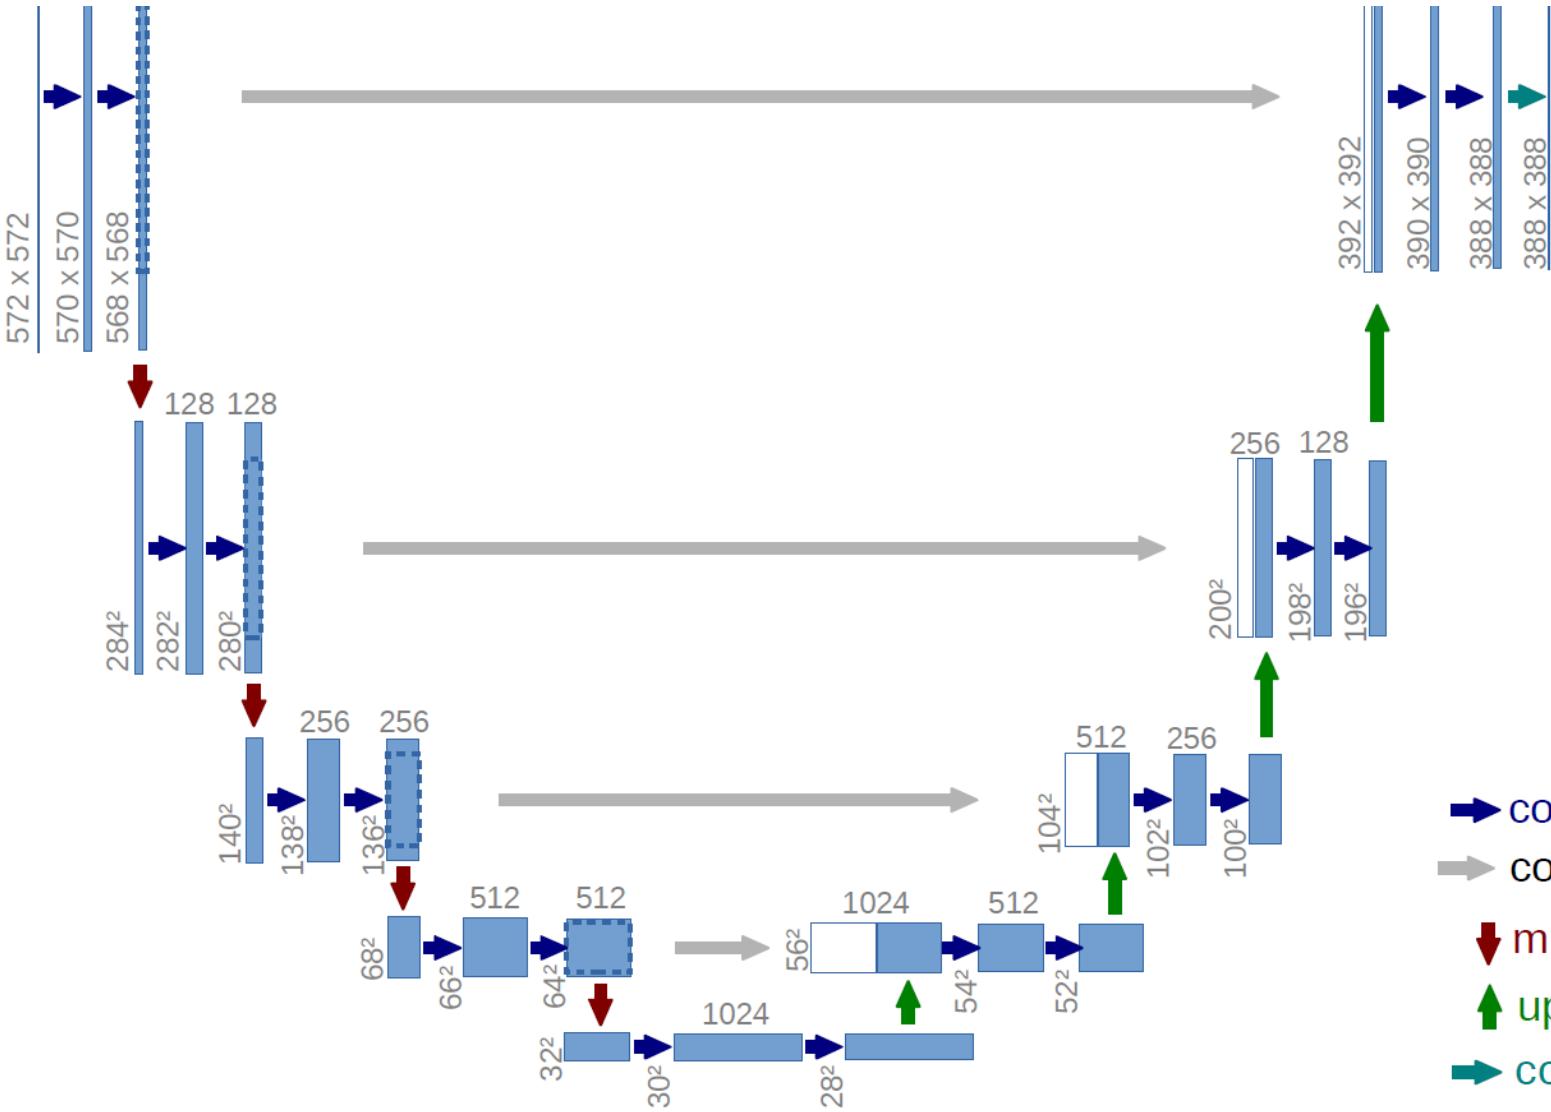




A small orange rectangular bar is located in the top-left corner of the slide.

UNet

And its applications



- conv 3x3, ReLU
- copy and crop
- ↓ max pool 2x2
- ↑ up-conv 2x2
- conv 1x1

U-Net applications

Image Segmentation

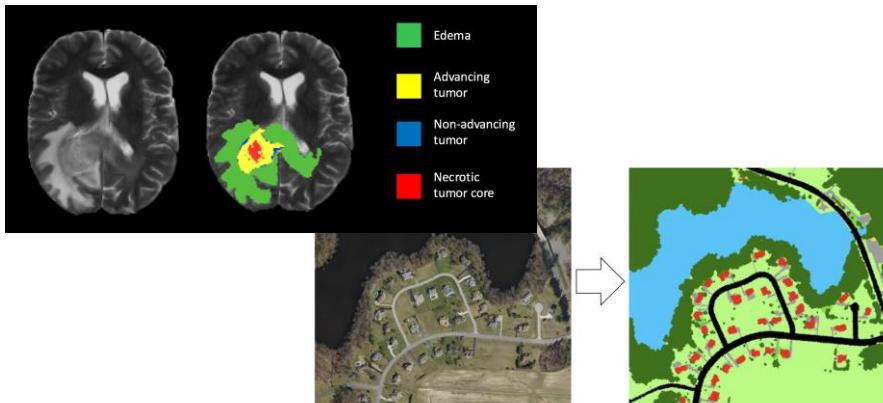


Image Denoising

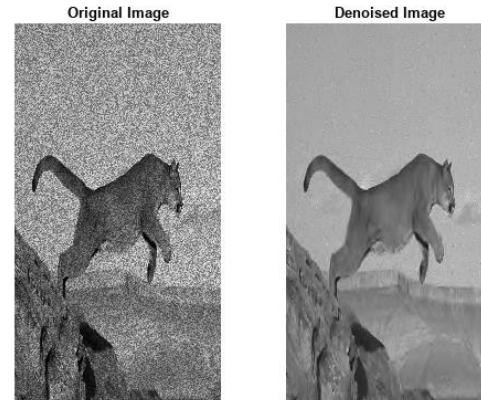


Image Upscaling

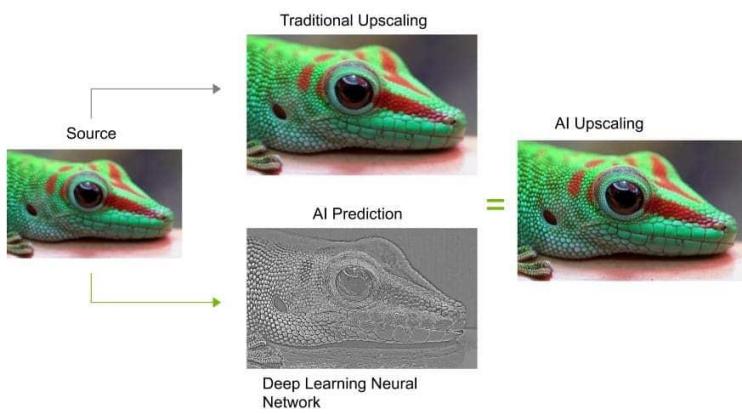


Image Generation

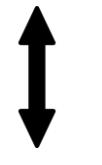


Data augmentation

Increased dataset diversity
Prevents overfitting
Enhanced feature learning



Horizontal Flip



Vertical Flip

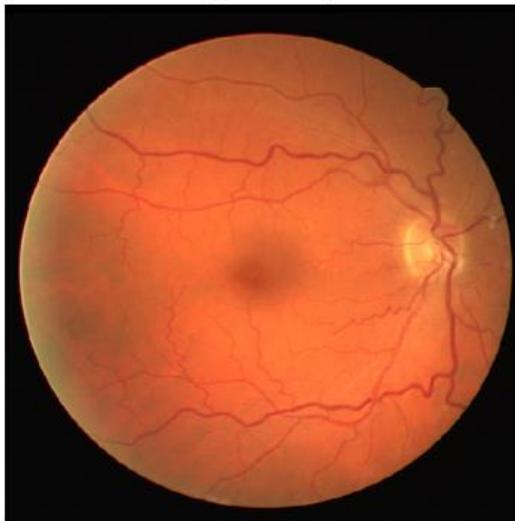


Rotation



Brightness Adjustment

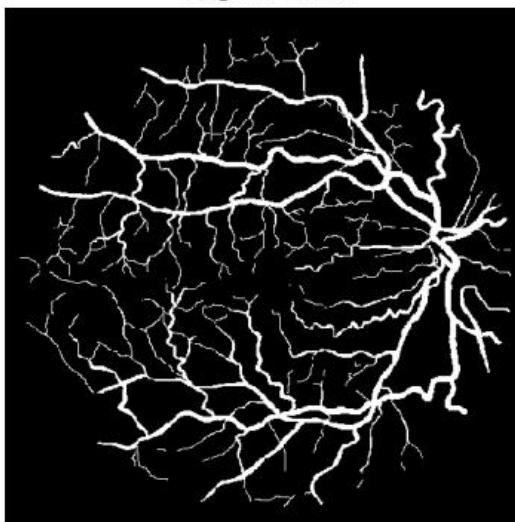
Original Image



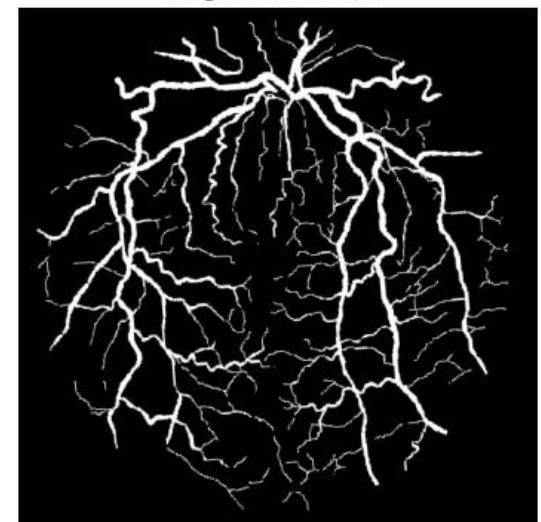
Augmented Image



Original Mask



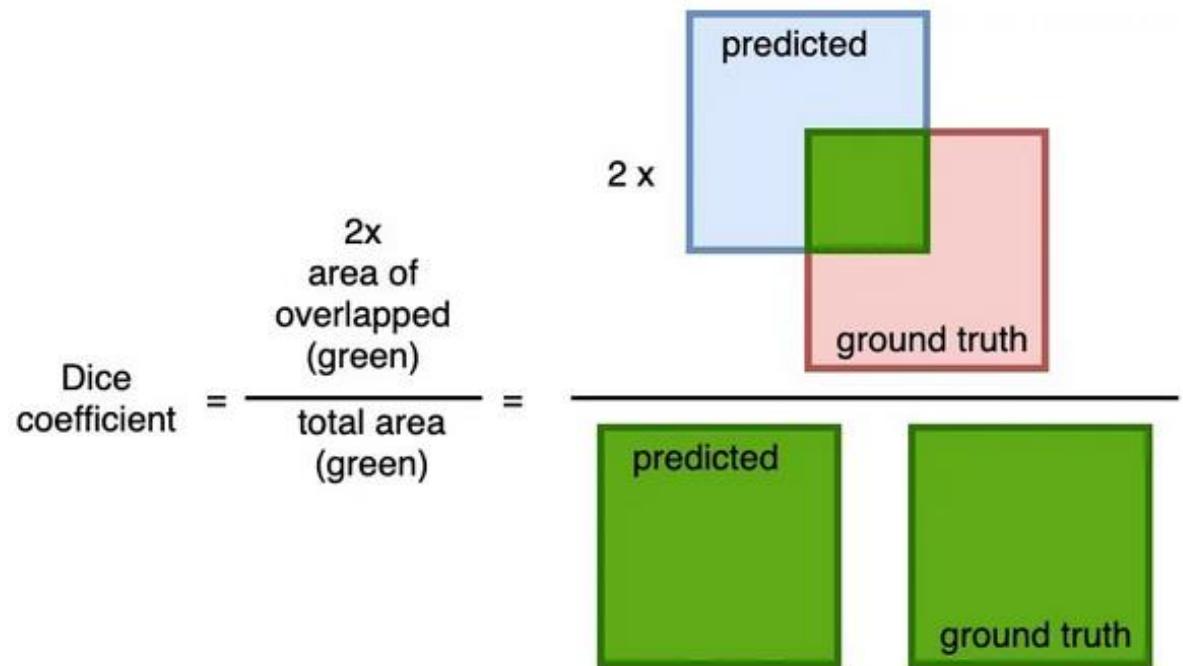
Augmented Mask



DICE loss

Solves imbalance problem by increasing weight of overlap
(high imbalance typical for medical imaging)

Quantification of the similarity between the two sets

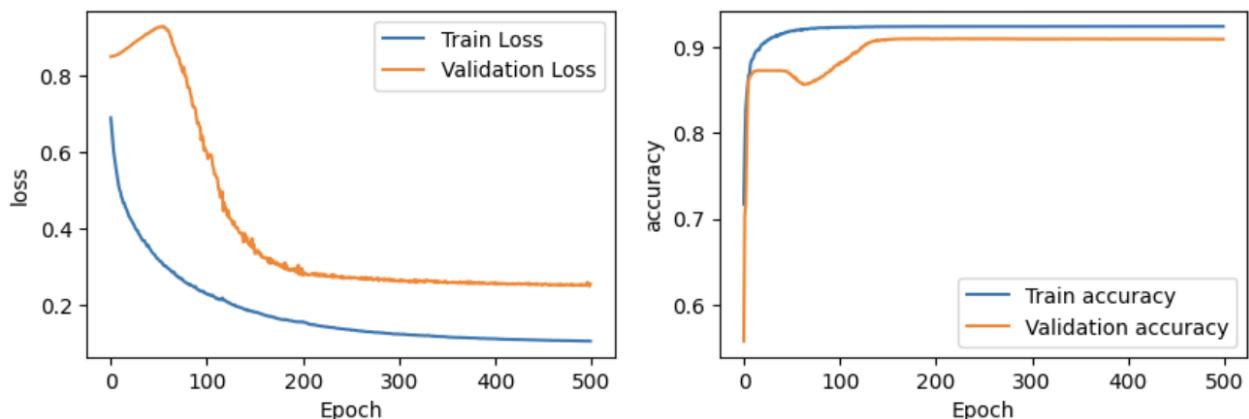
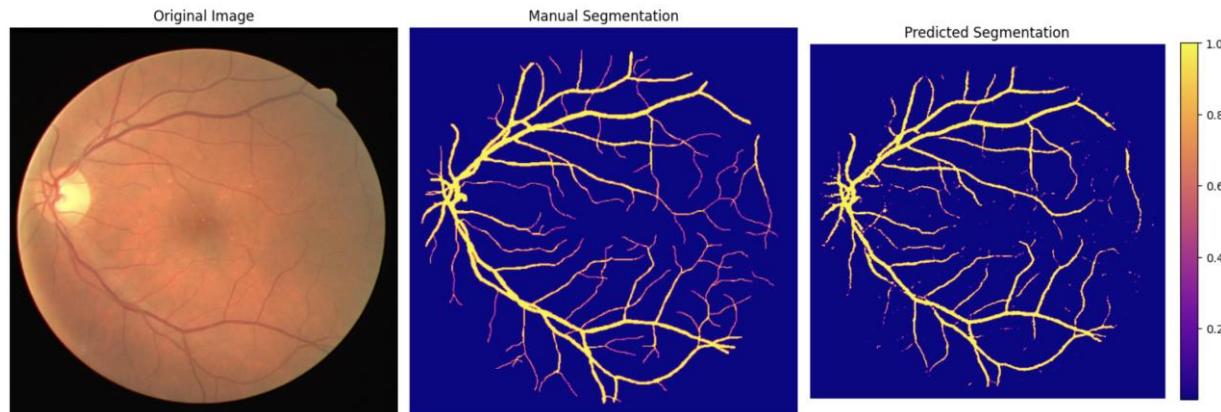


U-NET: Baseline Results

U-Net: 31,039,425 trainable parameters

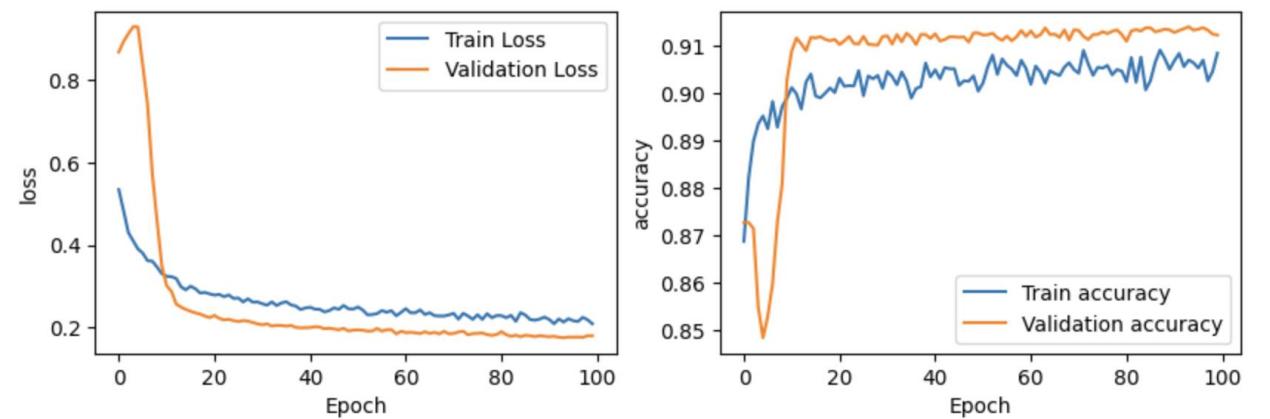
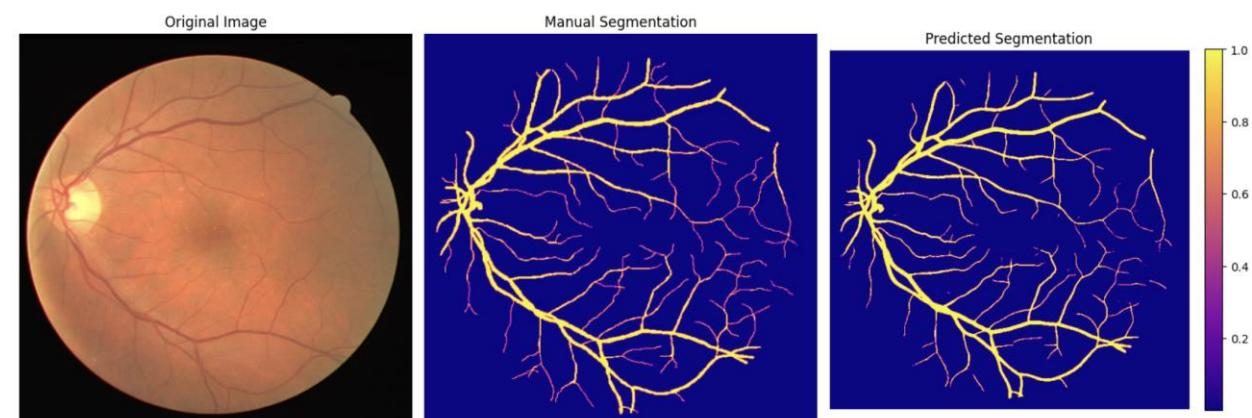
Original Dataset

- Validation DICE: 0,7499



Augmented Data

- Validation DICE: 0,8246

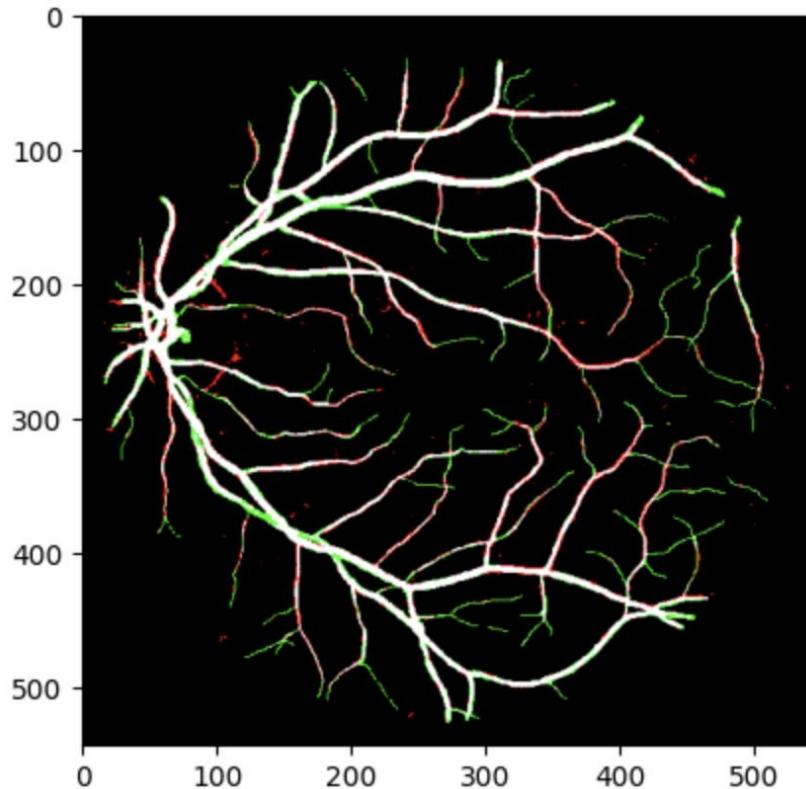


U-NET: Baseline Results

U-Net: 31,039,425 trainable parameters

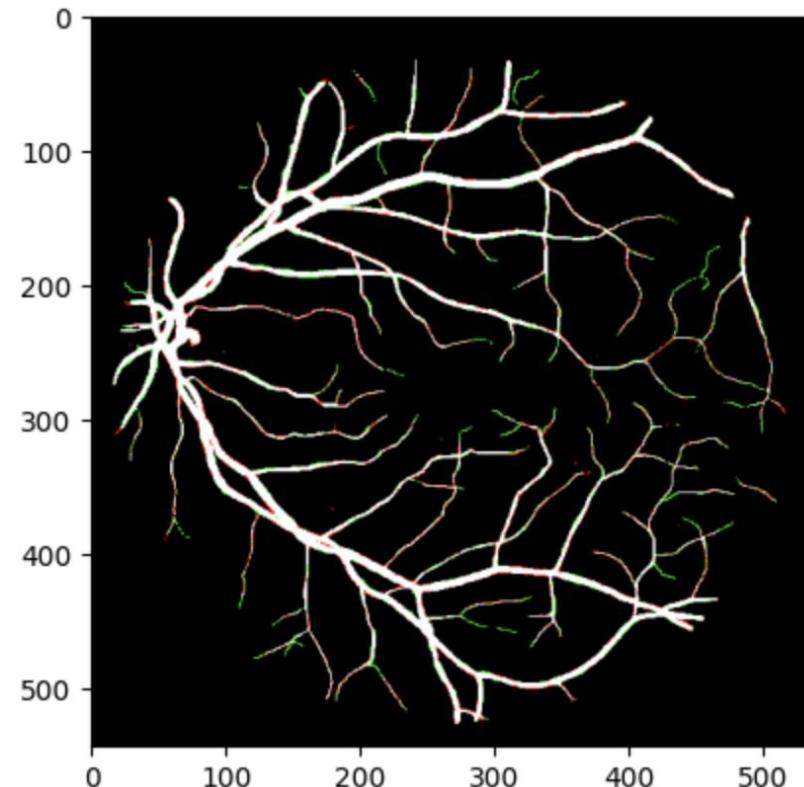
Original Dataset

- Validation DICE: 0,7499



Augmented Data

- Validation DICE: 0,8246

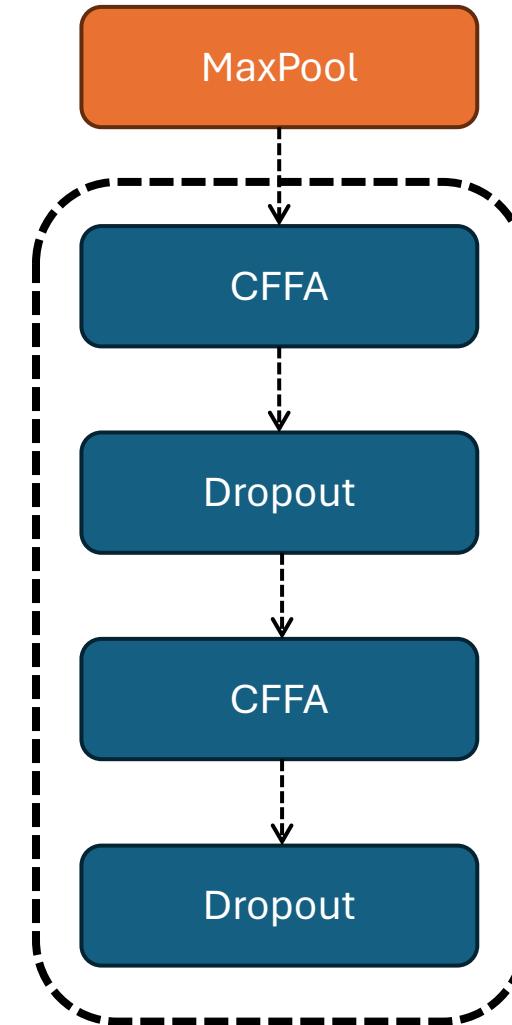
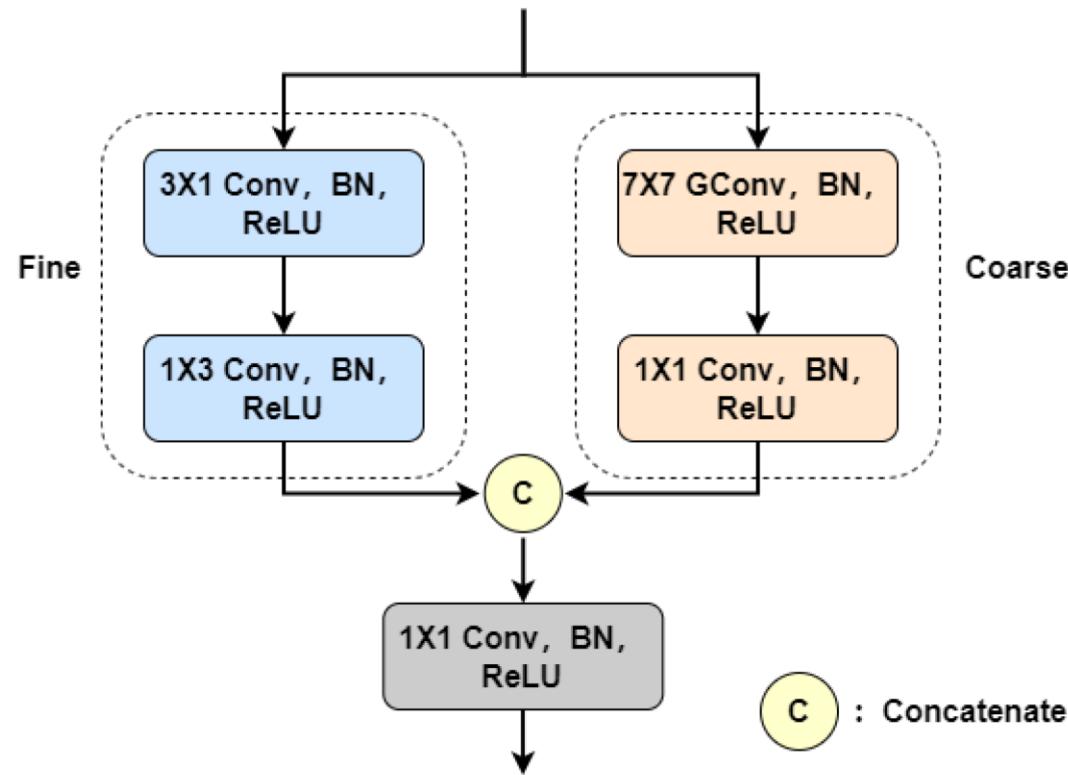


TP	
TN	
FP	
FN	

Advanced Architectures

- CBAM-UNet (With Coarse Fine Feature Aggregation)
 - CMP-UNet
 - Deformable UNet
-

Down: CFFA



```
def CFFA(input, num_filters):
    fine = Conv2D(num_filters, (3, 1), padding="same")(input)
    fine = BatchNormalization()(fine)
    fine = Activation("relu")(fine)

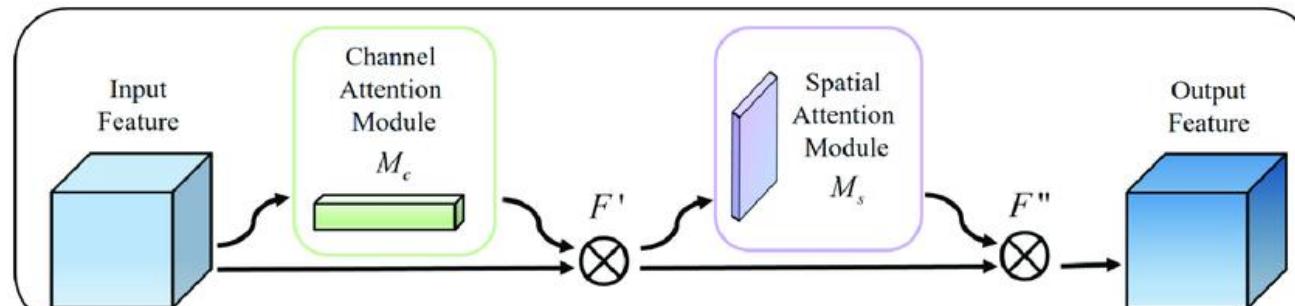
    fine = Conv2D(num_filters, (1, 3), padding="same")(fine)
    fine = BatchNormalization()(fine)
    fine = Activation("relu")(fine)

    coarse = Conv2D(num_filters, 7, padding="same")(input)
    coarse = BatchNormalization()(coarse)
    coarse = Activation("relu")(coarse)

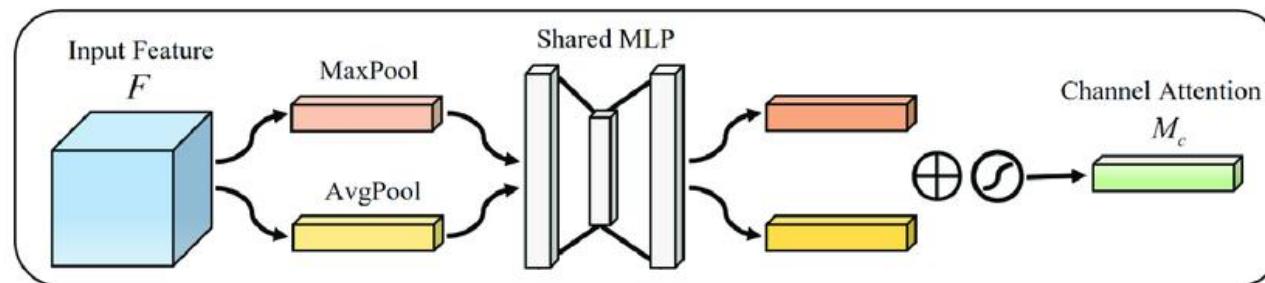
    coarse = Conv2D(num_filters, 1, padding="same")(coarse)
    coarse = BatchNormalization()(coarse)
    coarse = Activation("relu")(coarse)

    x = Concatenate()([fine, coarse])
    x = Conv2D(num_filters, 1, padding="same")(x)
    x = BatchNormalization()(x)
    x = Activation("relu")(x)
    return x
```

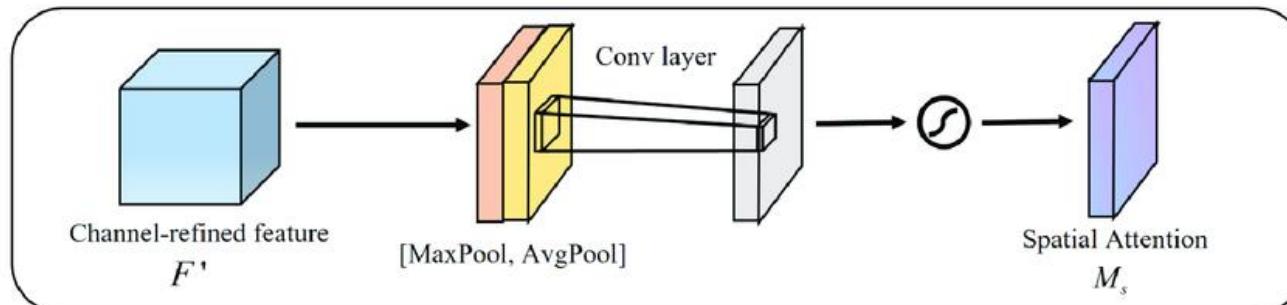
Bridge: Convolutional Block Attention Module



(a)



(b)



(c)

```
def channel_attention(input_feature, ratio=8):
    channel = input_feature.shape[-1]
    shared_layer_one = Dense(channel//ratio,
                             activation='relu',
                             kernel_initializer='he_normal',
                             use_bias=True,
                             bias_initializer='zeros')
    shared_layer_two = Dense(channel,
                             kernel_initializer='he_normal',
                             use_bias=True,
                             bias_initializer='zeros')
    avg_pool = GlobalAveragePooling2D()(input_feature)
    avg_pool = Reshape((1,1,channel))(avg_pool)
    assert avg_pool.shape[1:] == (1,1,channel)
    avg_pool = shared_layer_one(avg_pool)
    assert avg_pool.shape[1:] == (1,1,channel//ratio)
    avg_pool = shared_layer_two(avg_pool)
    assert avg_pool.shape[1:] == (1,1,channel)

    max_pool = GlobalMaxPooling2D()(input_feature)
    max_pool = Reshape((1,1,channel))(max_pool)
    assert max_pool.shape[1:] == (1,1,channel)
    max_pool = shared_layer_one(max_pool)
    assert max_pool.shape[1:] == (1,1,channel//ratio)
    max_pool = shared_layer_two(max_pool)
    assert max_pool.shape[1:] == (1,1,channel)

    cbam_feature = Add()([avg_pool,max_pool])
    cbam_feature = Activation('sigmoid')(cbam_feature)
    return Multiply()([input_feature, cbam_feature])
```

```
def spatial_attention(input_feature):
    kernel_size = 7
    channel = input_feature.shape[-1]
    cbam_feature = input_feature
    output_shape = (input_feature.shape[1], input_feature.shape[2], 1)

    avg_pool = Lambda(lambda x: tf.keras.backend.mean(x, axis=3, keepdims=True), output_shape)(cbam_feature)
    assert avg_pool.shape[-1] == 1
    max_pool = Lambda(lambda x: tf.keras.backend.max(x, axis=3, keepdims=True), output_shape)(cbam_feature)
    assert max_pool.shape[-1] == 1
    concat = Concatenate(axis=3)([avg_pool, max_pool])
    assert concat.shape[-1] == 2
    cbam_feature = Conv2D(filters=1,
                          kernel_size=kernel_size,
                          strides=1,
                          padding='same',
                          activation='sigmoid',
                          kernel_initializer='he_normal',
                          use_bias=False)(concat)
    assert cbam_feature.shape[-1] == 1

    return Multiply()([input_feature, cbam_feature])
```

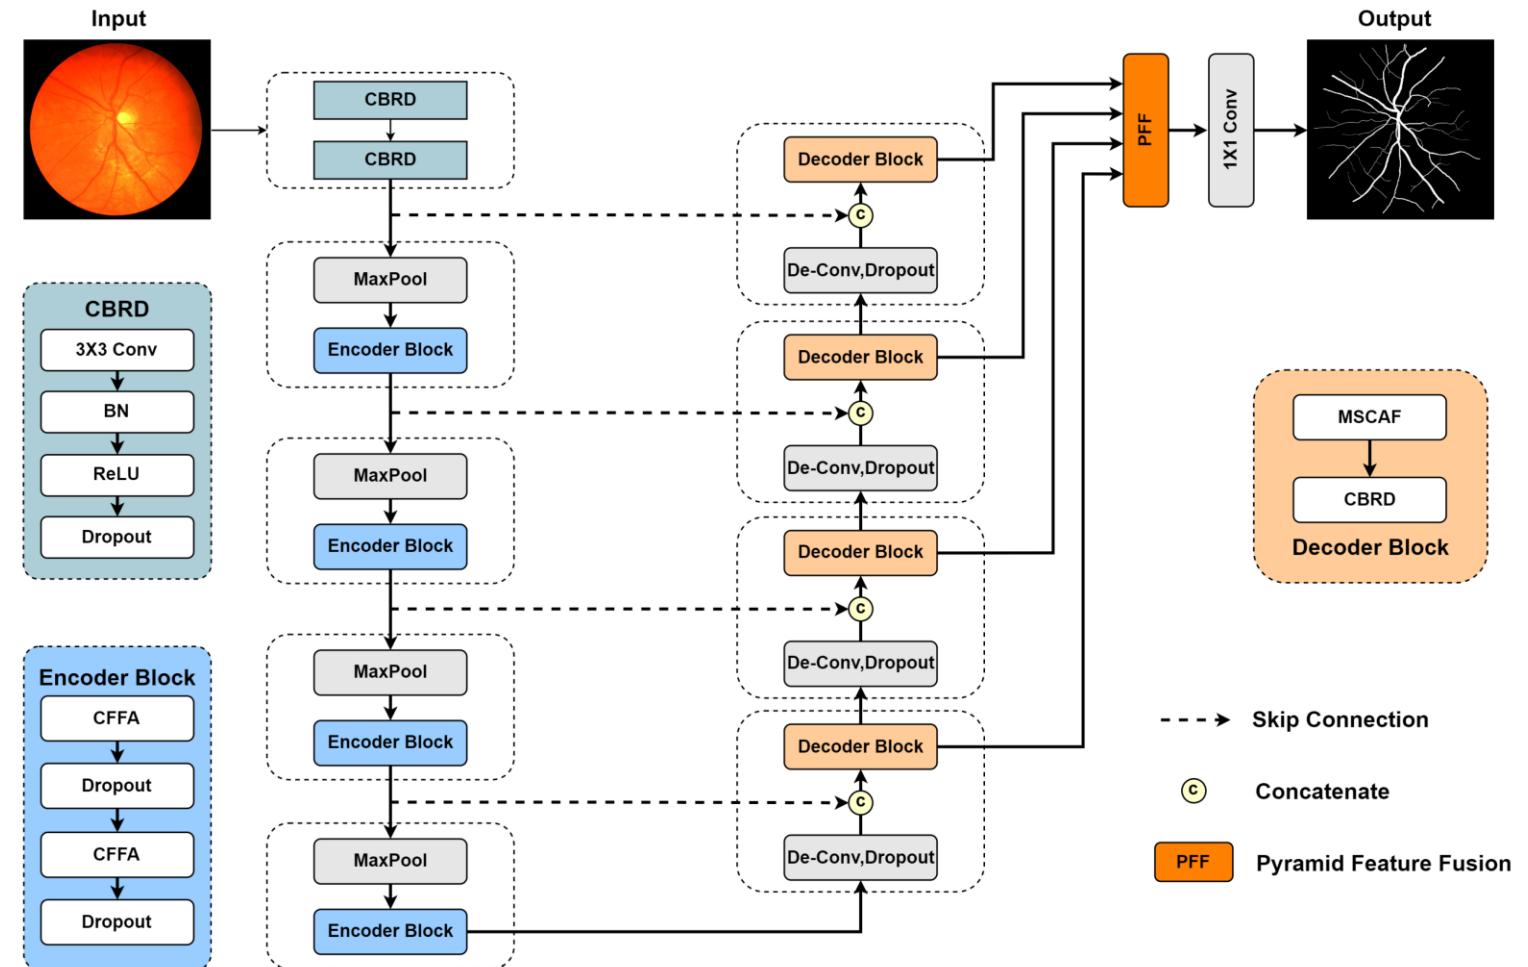


Architectures

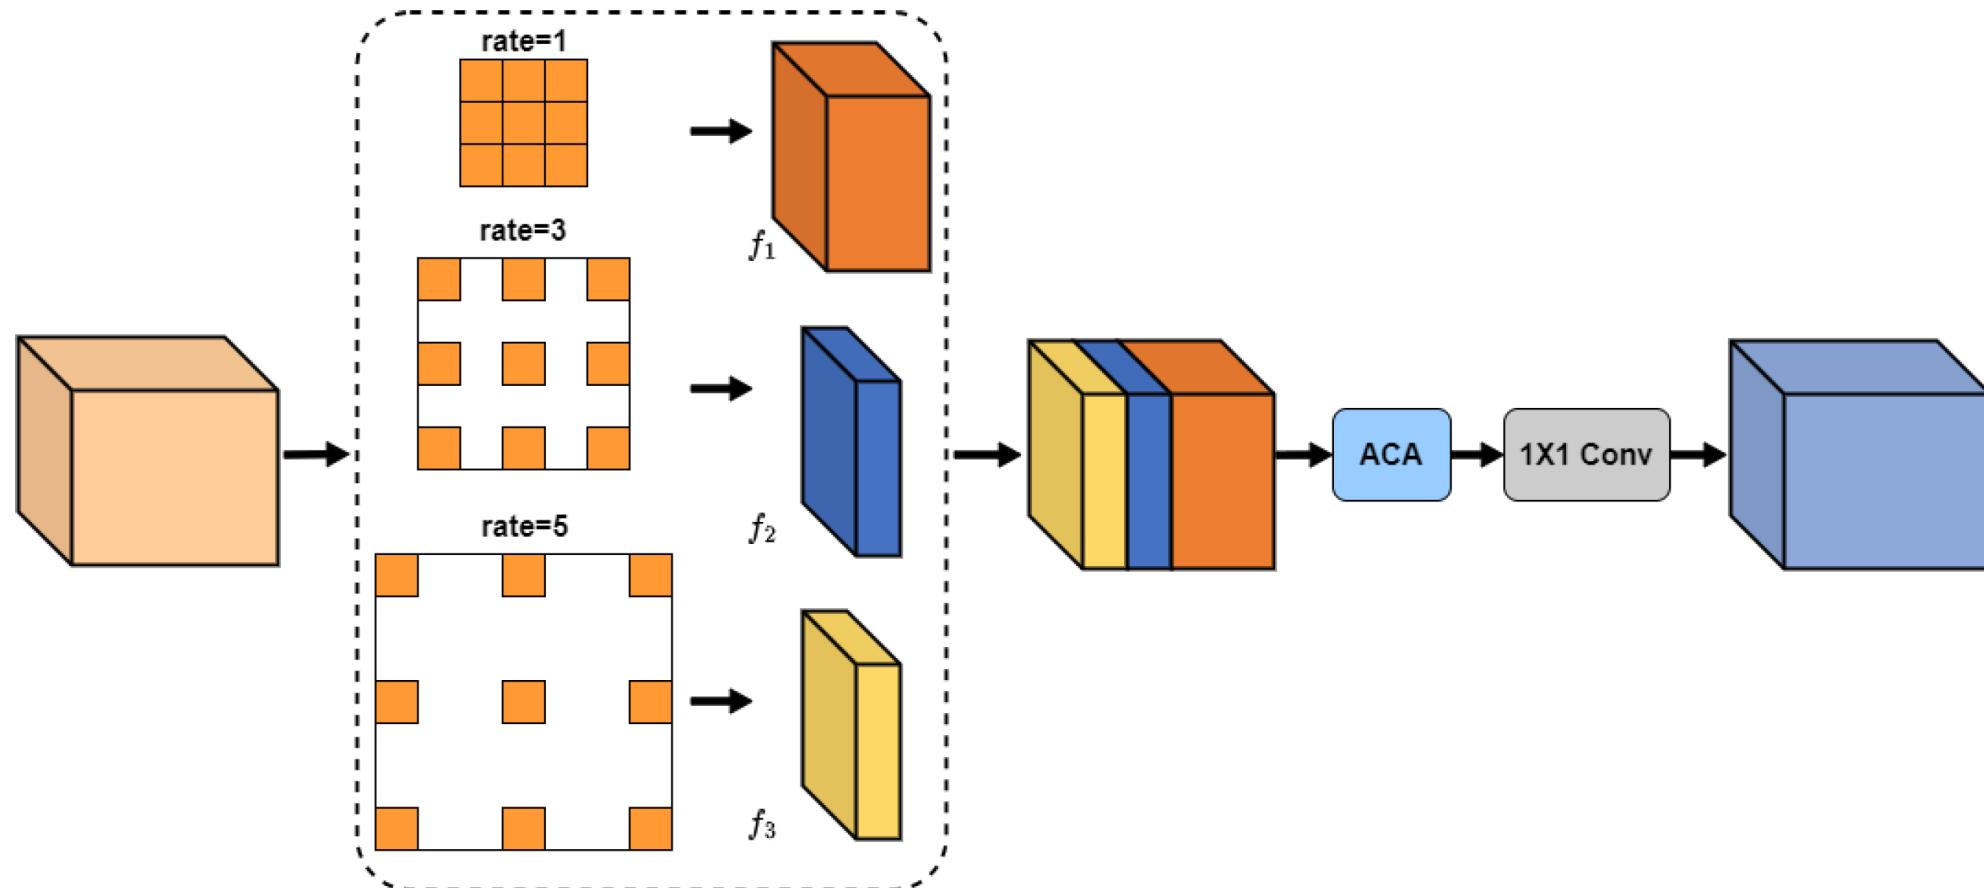
- CBAM-UNet (With Coarse Fine Feature Aggregation)
 - CMP-Unet
 - Deformable UNet
-

CMP-Unet

Gu Y., et al.



Multi-Scale Channel Adaptive Fusion



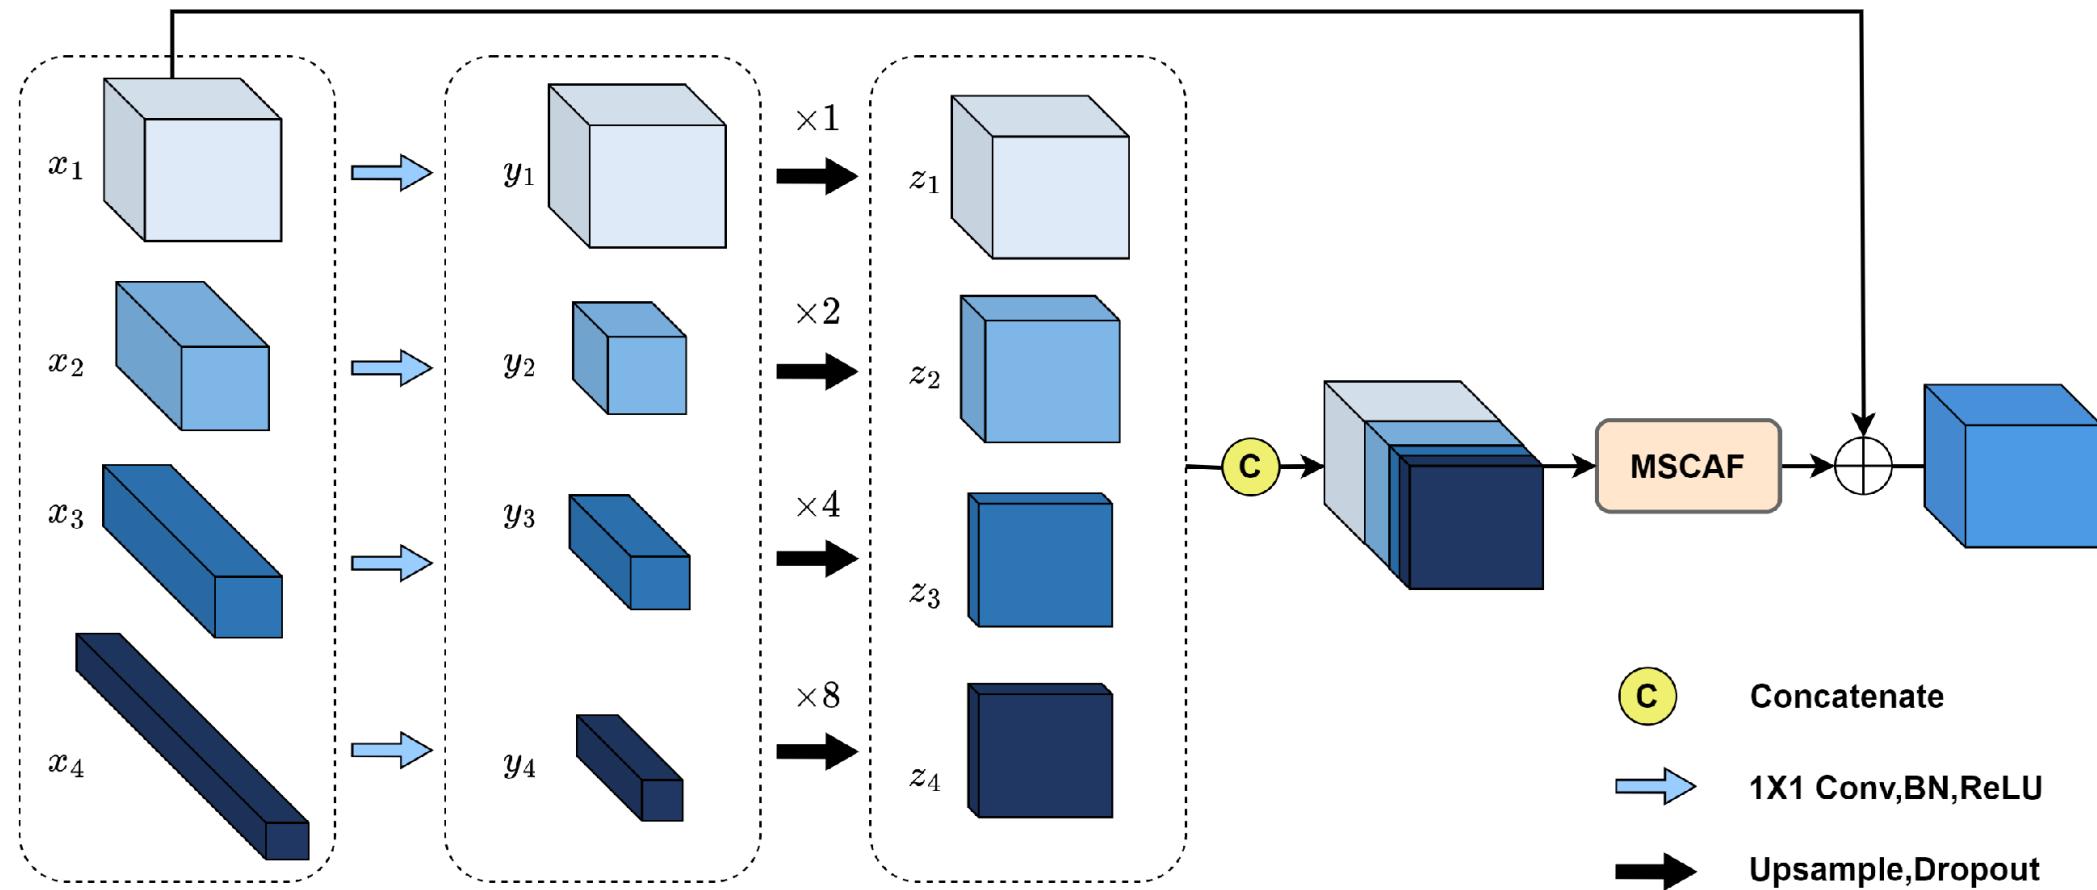
```
def MSCAF(input, out_channels, dropout):
    x1 = Conv2D(filters=out_channels//2, kernel_size=3, padding='same', use_bias=False)(input)
    x1 = BatchNormalization()(x1)
    x1 = Activation("relu")(x1)

    x2 = Conv2D(filters=out_channels//4, kernel_size=3, dilation_rate=3, padding='same', use_bias=False)(input)
    x2 = BatchNormalization()(x2)
    x2 = Activation("relu")(x2)

    x3 = Conv2D(filters=out_channels//4, kernel_size=3, dilation_rate=5, padding='same', use_bias=False)(input)
    x3 = BatchNormalization()(x3)
    x3 = Activation("relu")(x3)

    out = Concatenate()([x1, x2, x3])
    out = channel_attention(out)
    out = Conv2D(out_channels, 1, use_bias = False)(out)
    out = BatchNormalization()(out)
    out = Activation("relu")(out)
    out = Dropout(dropout)(out)
    return(out)
```

Pyramid Feature Fusion



```
def PFF(input1, input2, input3, input4, num_filters):
    x1 = Conv2D(num_filters // 8, 1, use_bias = False)(input1)
    x1 = BatchNormalization()(x1)
    x1 = Activation("relu")(x1)
    x1 = UpSampling2D(8, interpolation='bilinear')(x1)

    x2 = Conv2D(num_filters // 8, 1, use_bias = False)(input2)
    x2 = BatchNormalization()(x2)
    x2 = Activation("relu")(x2)
    x2 = UpSampling2D(4, interpolation='bilinear')(x2)

    x3 = Conv2D(num_filters // 4, 1, use_bias = False)(input3)
    x3 = BatchNormalization()(x3)
    x3 = Activation("relu")(x3)
    x3 = UpSampling2D(2, interpolation='bilinear')(x3)

    x4 = Conv2D(num_filters // 2, 1, use_bias = False)(input4)
    x4 = BatchNormalization()(x4)
    x4 = Activation("relu")(x4)

    out = Concatenate()([x1, x2, x3, x4])
    print(out.shape)
    out = MSCAF(out, num_filters) + input4
    return out
```

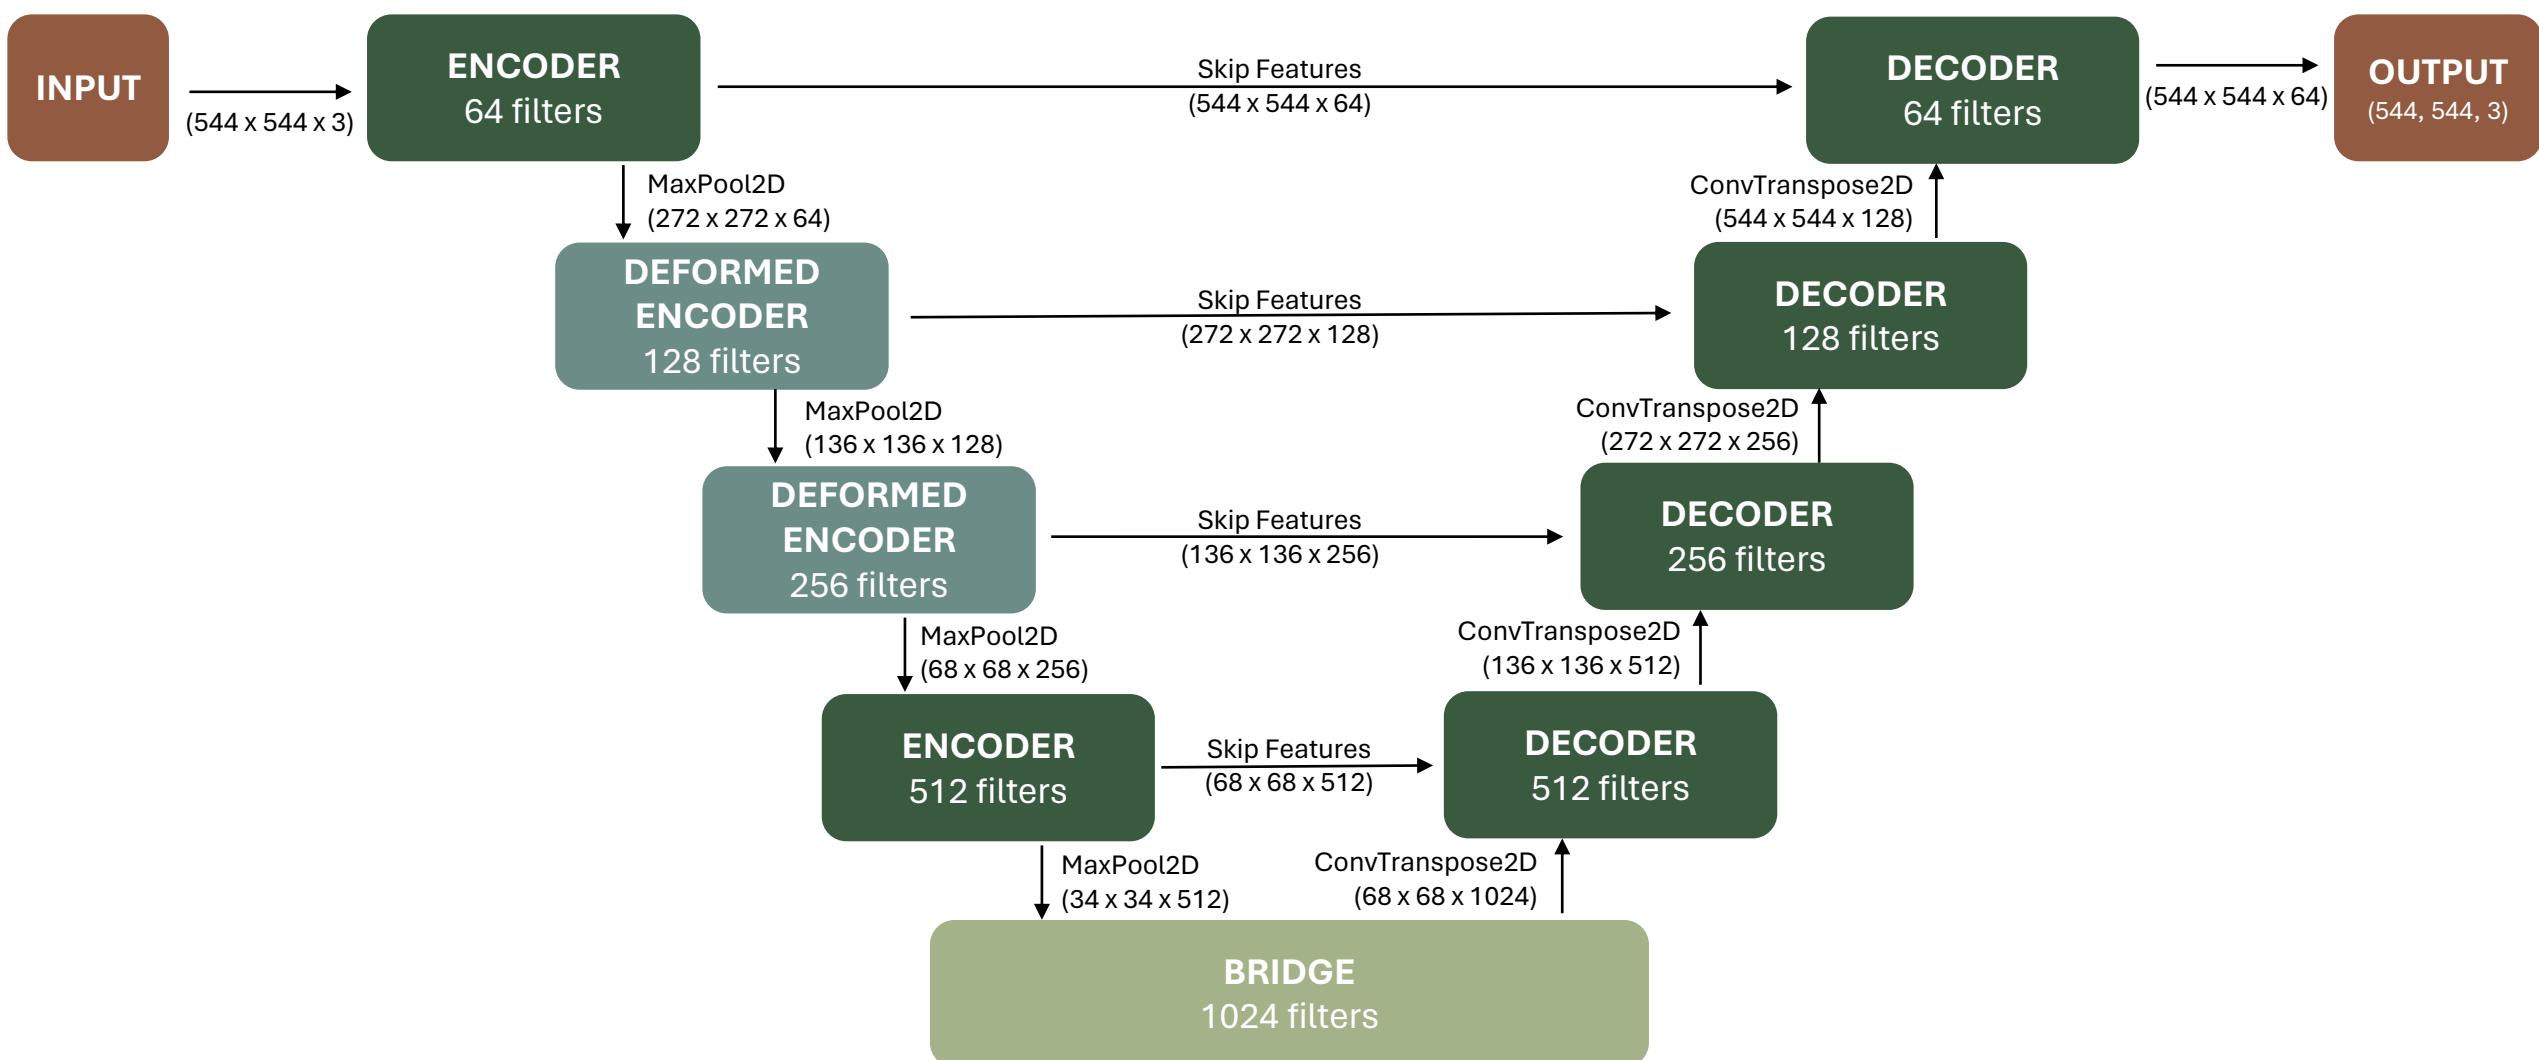
Architectures

- CBAM-UNet (With Coarse Fine Feature Aggregation)
 - CMP-UNet
 - Deformable UNet
-

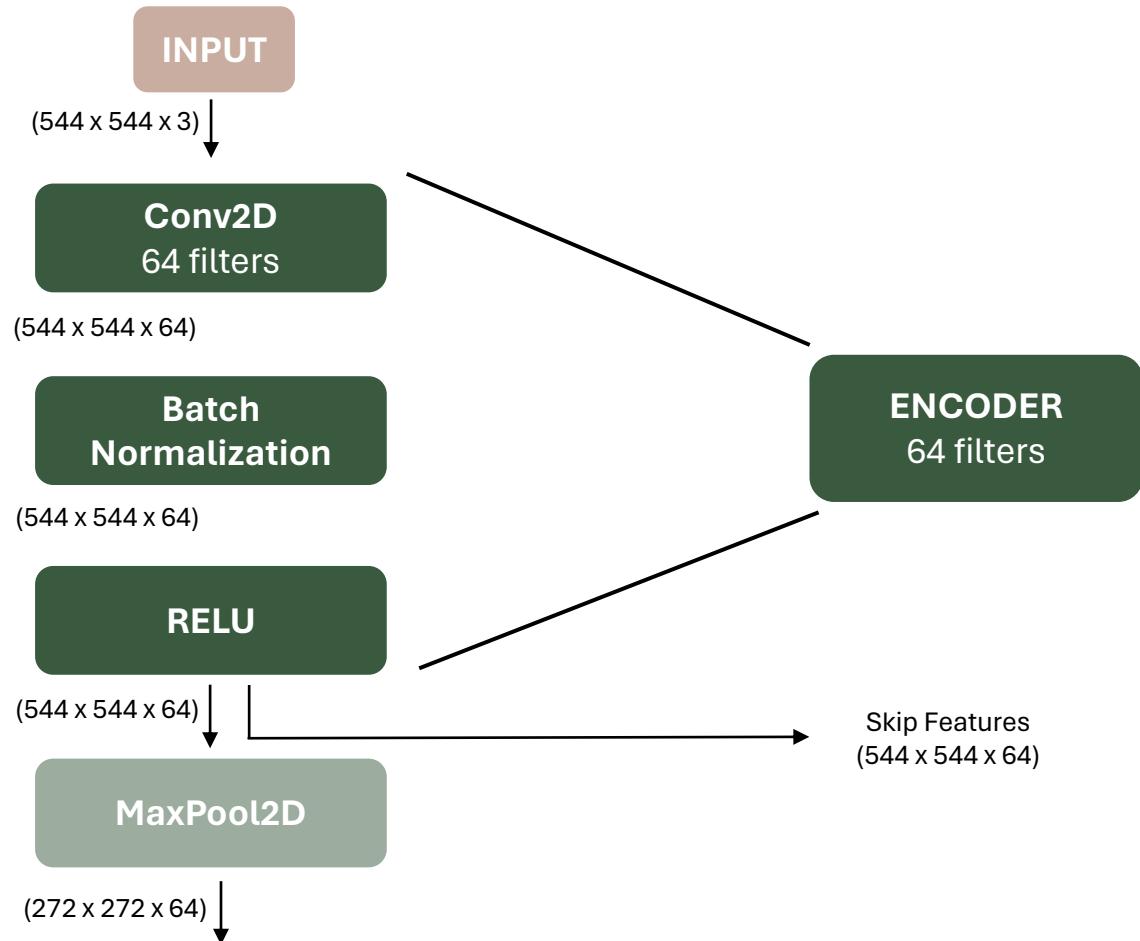
DU-Net

Jin, Q. et al.

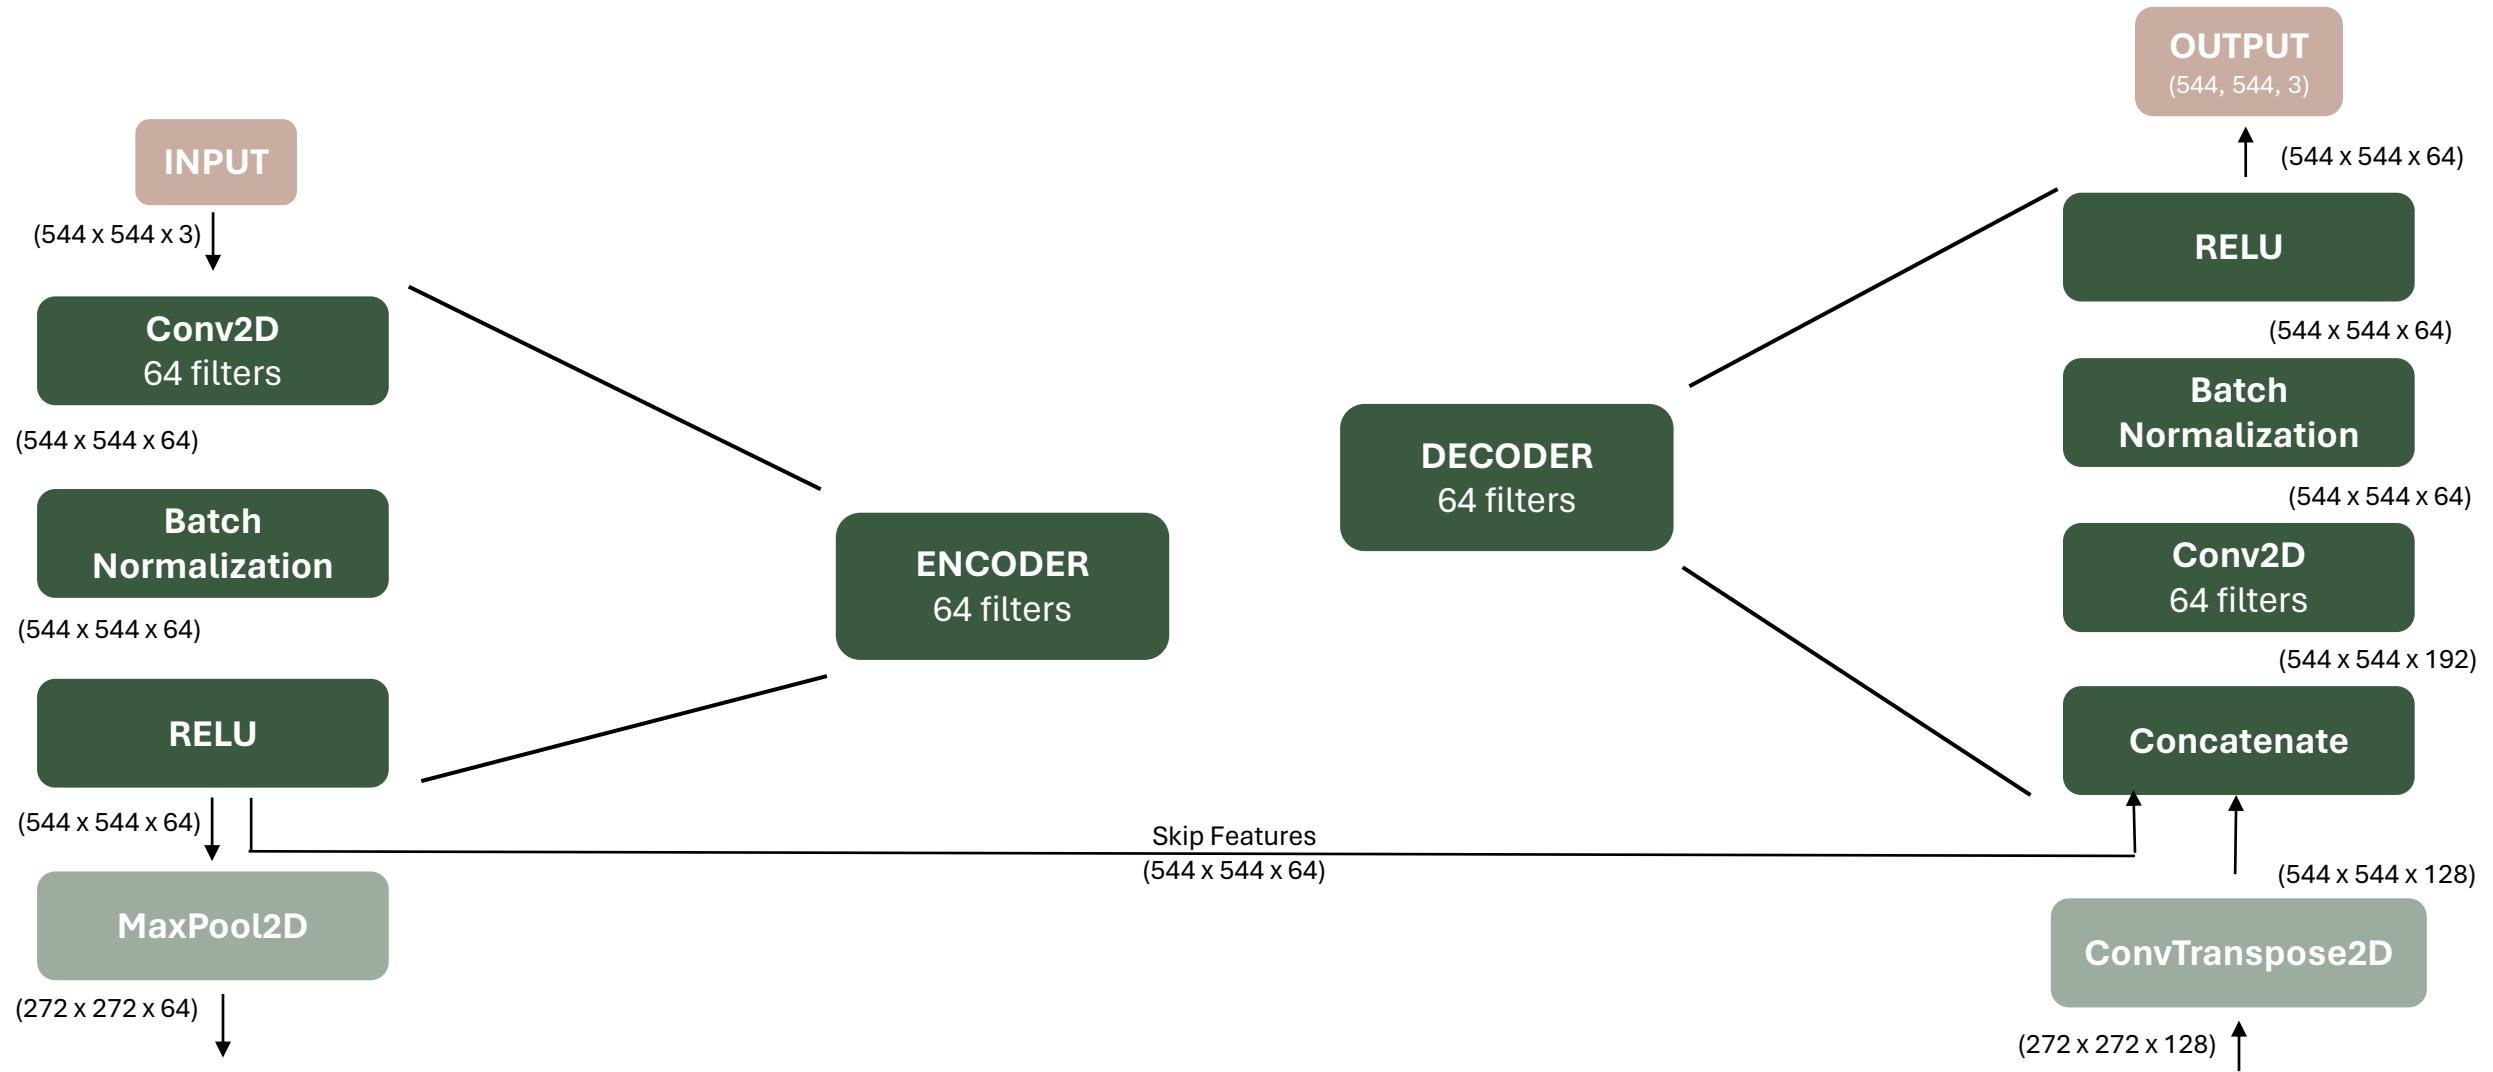
Architecture



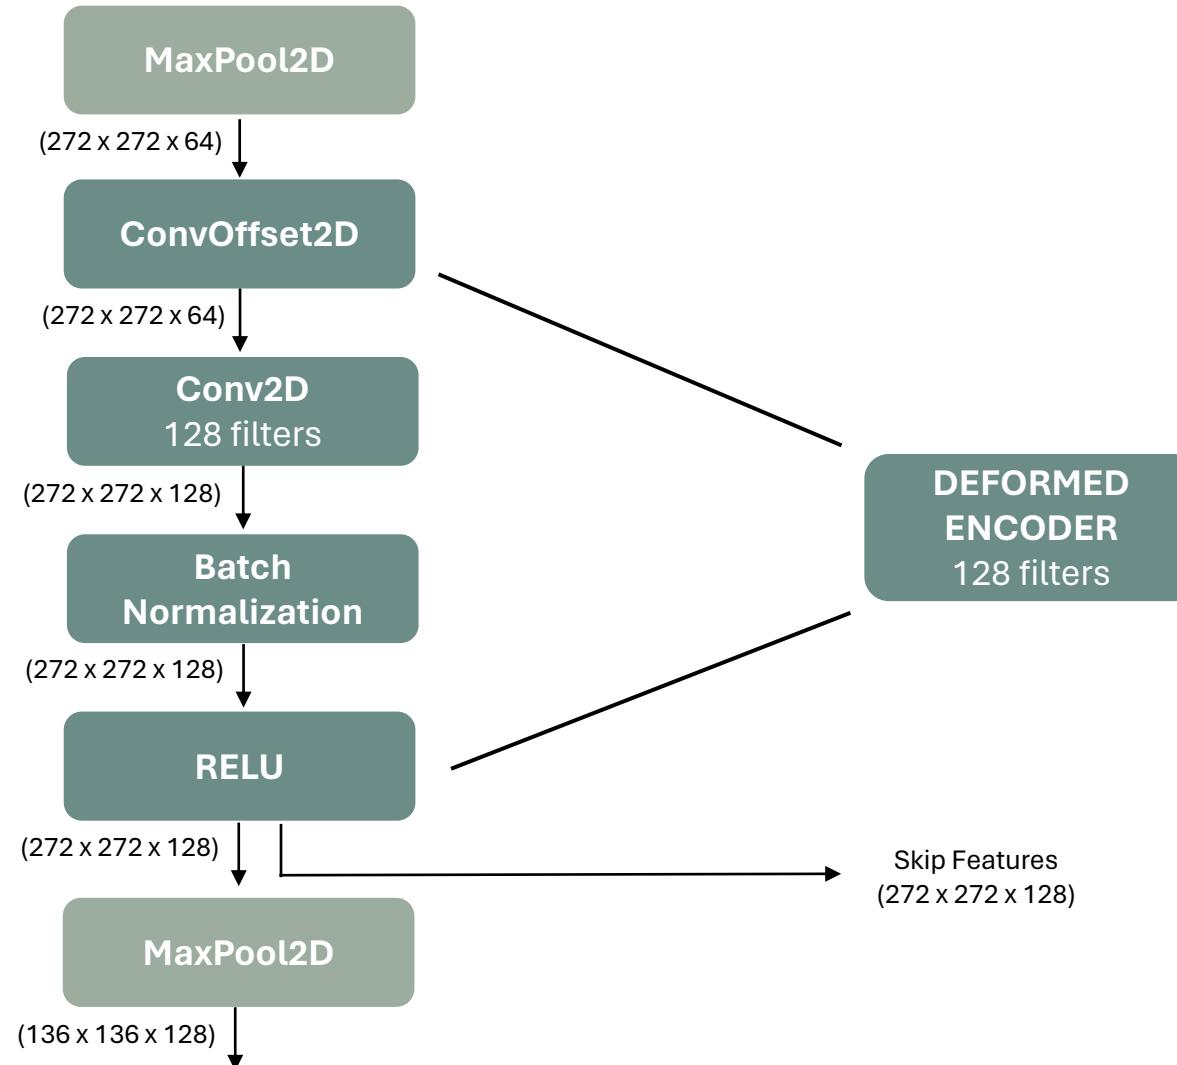
Architecture



Architecture



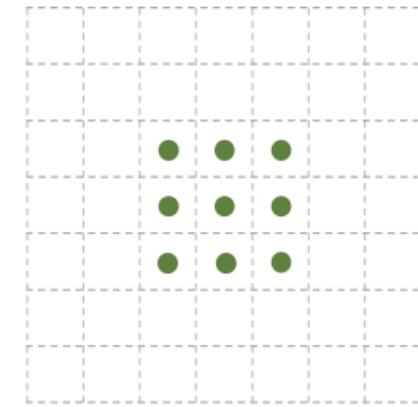
Architecture



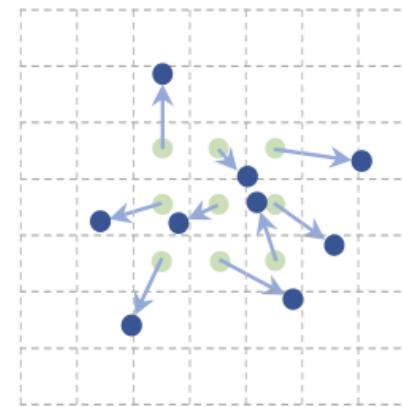
Architecture

ConvOffset2D

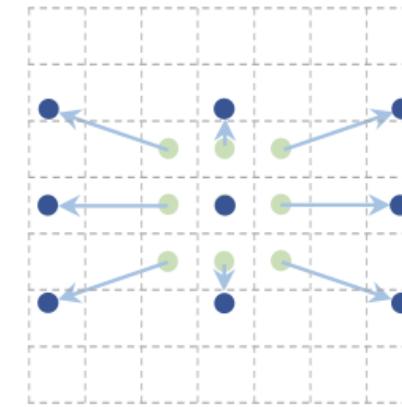
Returns deformed
feature map



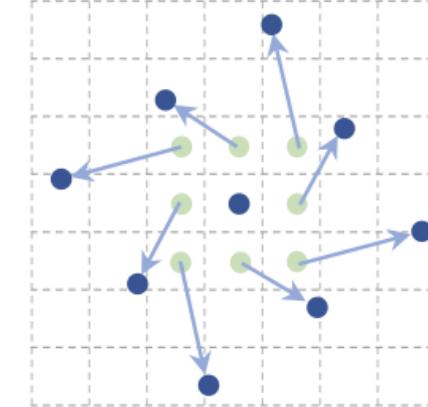
(a)



(b)



(c)



(d)

Image source:
Deformable Convolutional Networks – Dai, J. et al.

Architecture

ConvOffset2D

Returns deformed
feature map

- Calculates x and y offset for each "pixel" in each feature map

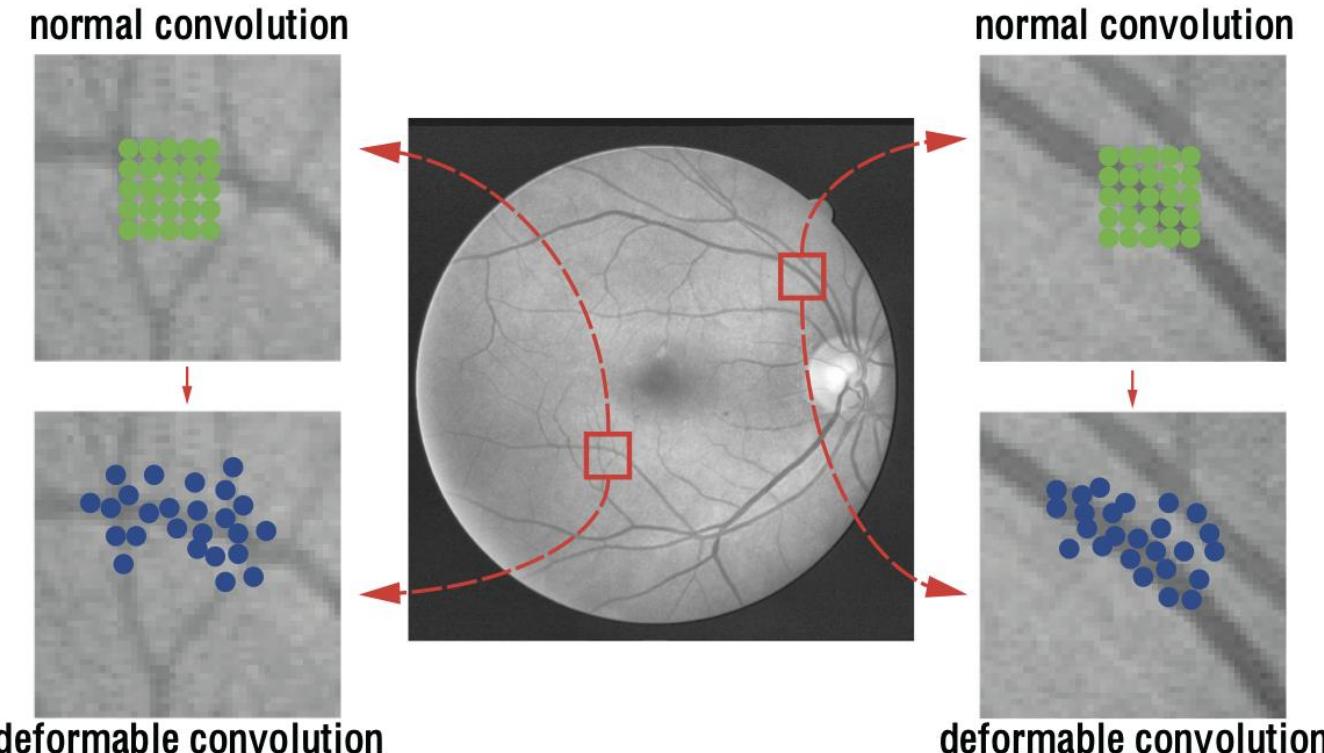


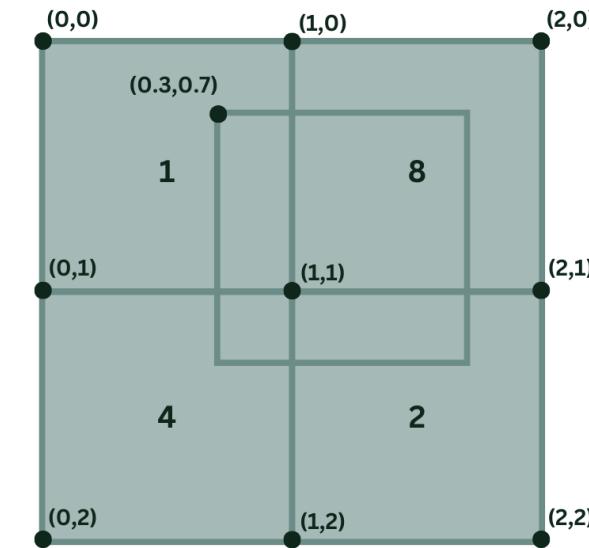
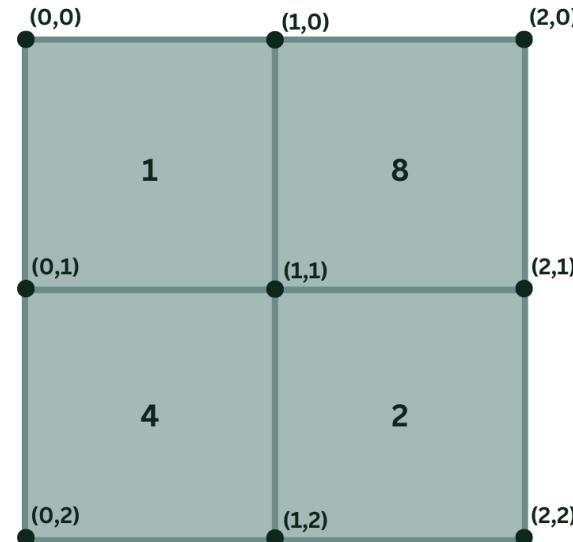
Image source:
DUNet: A deformable network for retinal vessel segmentation - Jin, Q. et al.

Architecture

ConvOffset2D

- Calculates x and y offset (Δp_n) for each "pixel" in each feature map
- In the deformed feature map, value $x(p_0 + p_n)$ is replaced by $x(p_0 + p_n + \Delta p_n)$.

Bilinear Interpolation



$$x(p) = \sum_q (g(q_x + p_x) \cdot g(q_y + p_y) \cdot x(p))$$

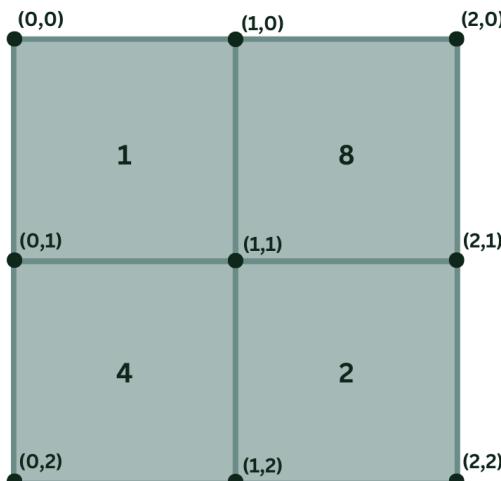
with $g(a, b) = \max(0, 1 - |a - b|)$

Architecture

ConvOffset2D

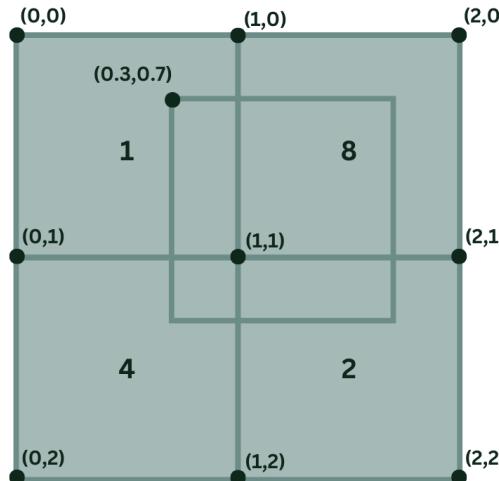
- Calculates x and y offset (Δp_n) for each "pixel" in each feature map
- In the deformed feature map, value $x(p_0 + p_n)$ is replaced by $x(p_0 + p_n + \Delta p_n)$.

Bilinear Interpolation



$$x(p) = \sum_q (g(q_x, p_x) \cdot g(q_y, p_y) \cdot x(p))$$

with $g(a, b) = \max(0, 1 - |a - b|)$



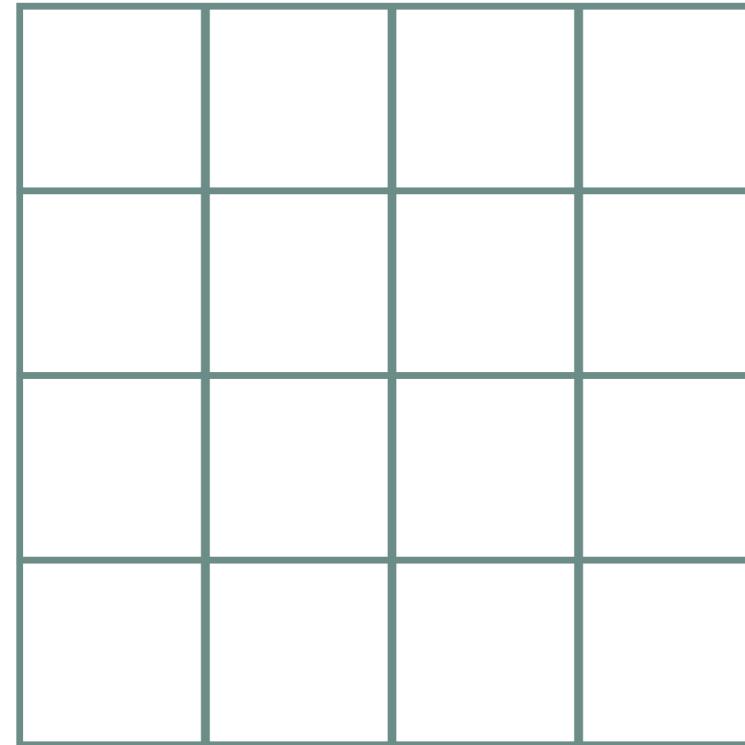
$$\begin{aligned} x(p) &= \max(0, 1 - |0 - 0.7|) \cdot \max(0, 1 - |0 - 0.3|) \cdot x(0,0) \\ &\quad + \max(0, 1 - |1 - 0.7|) \cdot \max(0, 1 - |0 - 0.3|) \cdot x(1,0) \\ &\quad + \max(0, 1 - |0 - 0.7|) \cdot \max(0, 1 - |1 - 0.3|) \cdot x(0,1) \\ &\quad + \max(0, 1 - |1 - 0.7|) \cdot \max(0, 1 - |1 - 0.3|) \cdot x(1,1) \end{aligned}$$

$$\begin{aligned} x(p) &= 0.3 \cdot 0.7 \cdot 1 \\ &\quad + 0.7 \cdot 0.7 \cdot 8 \\ &\quad + 0.3 \cdot 0.3 \cdot 4 \\ &\quad + 0.7 \cdot 0.3 \cdot 2 \end{aligned}$$

$$x(p) = 4.91 \approx 5$$

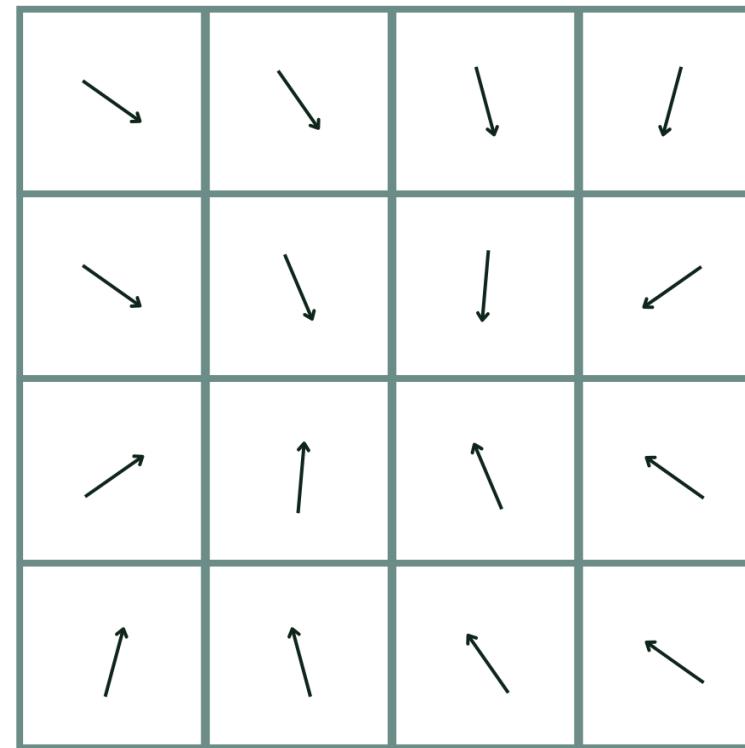
Architecture – Implementation Layer

ConvOffset2D



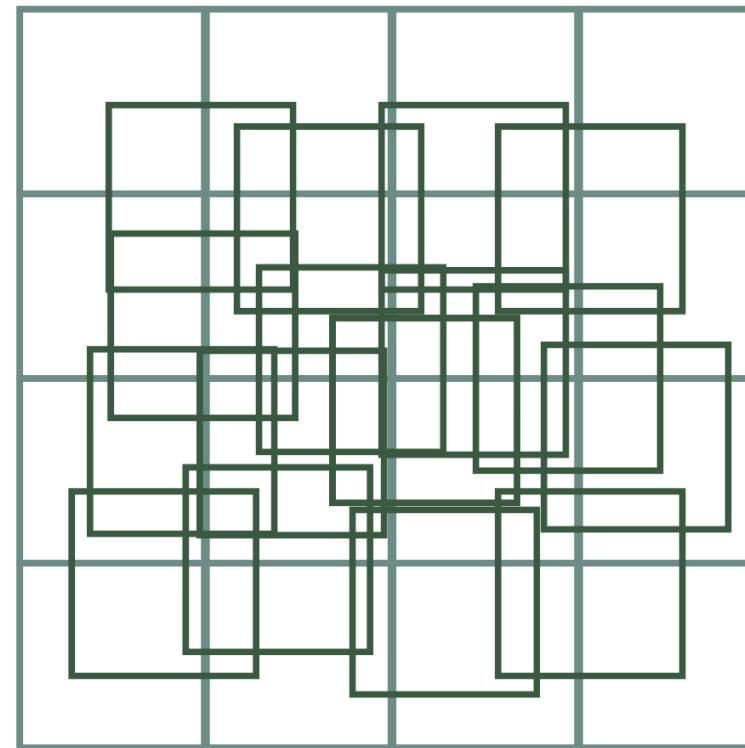
Architecture – Implementation Layer

ConvOffset2D



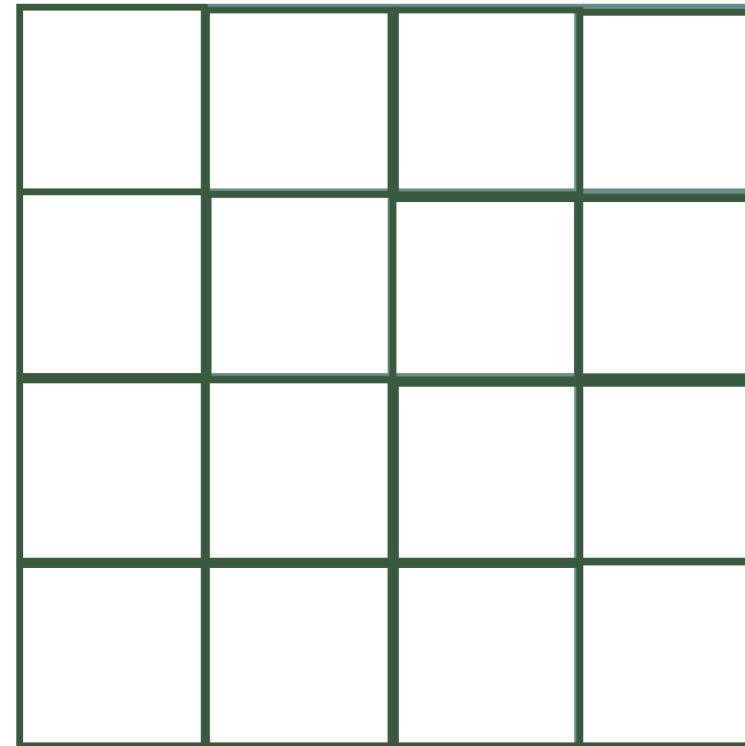
Architecture – Implementation Layer

ConvOffset2D



Architecture – Implementation Layer

ConvOffset2D



Architecture – Implementation Layer

ConvOffset2D

```
10 usages
class ConvOffset2D(Conv2D):
    def __init__(self, filters, init_normal_stddev=0.01, **kwargs):
        self.filters = filters
        super(ConvOffset2D, self).__init__(
            self.filters * 2, (3, 3), padding='same', use_bias=False,
            kernel_initializer=RandomNormal(0, init_normal_stddev),
            **kwargs
        )
        self.x_offsets = None

    def call(self, x):
        x_shape = x.get_shape()
        offsets = super(ConvOffset2D, self).call(x) # (b, h, w, 2c)

        # offsets: (b*c, h, w, 2)
        offsets = self._to_bc_h_w_2(offsets, x_shape)

        # x: (b*c, h, w)
        x = self._to_bc_h_w(x, x_shape)

        # X_offset: (b*c, h, w)
        x_offset = tf_batch_map_offsets(x, offsets)

        # x_offset: (b, h, w, c)
        x_offset = self._to_b_h_w_c(x_offset, x_shape)

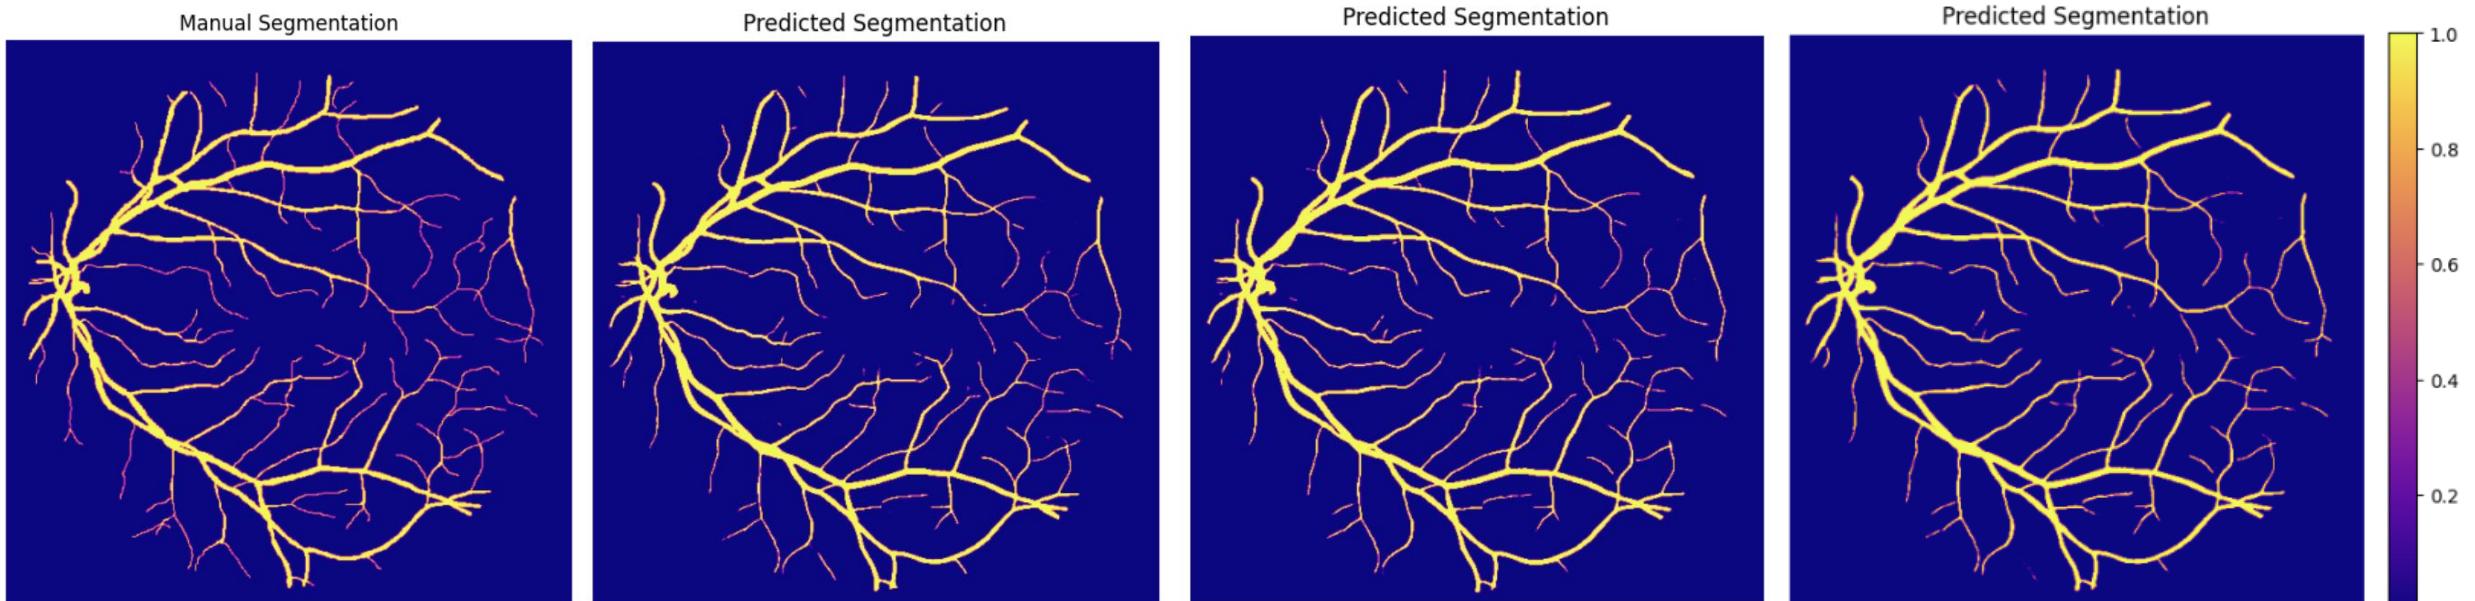
        return x_offset
```

Based on:

<https://github.com/kastnerkyle/deform-conv/tree/master>

Performance

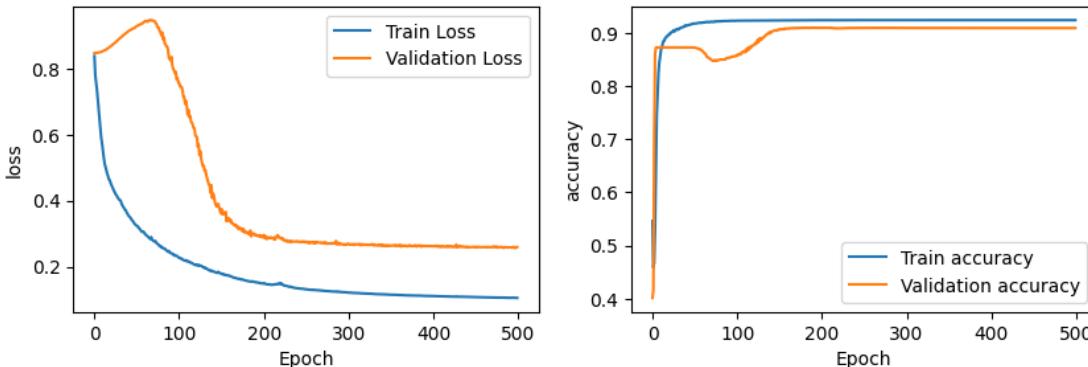
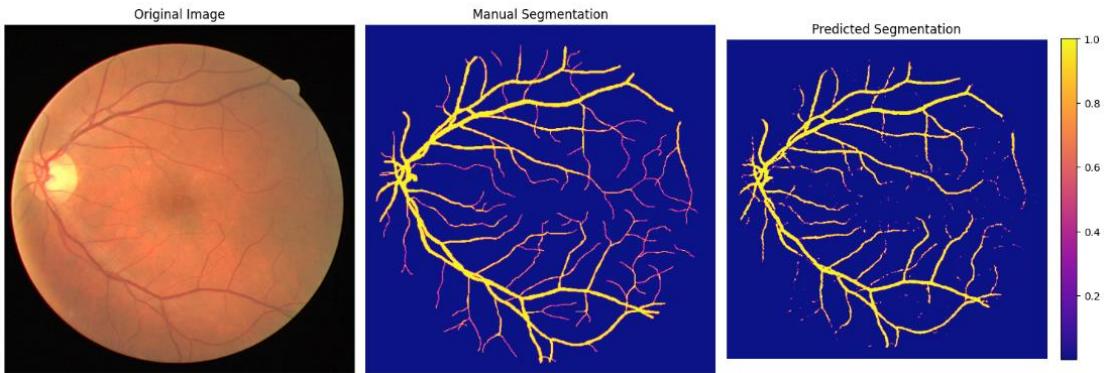
	U-Net	Deformable Encoder	DU-Net
Validation DICE	0.8246	0.8243	0.8233
Trainable Parameters	31 039 425	32 886 721	34 729 921



Performance

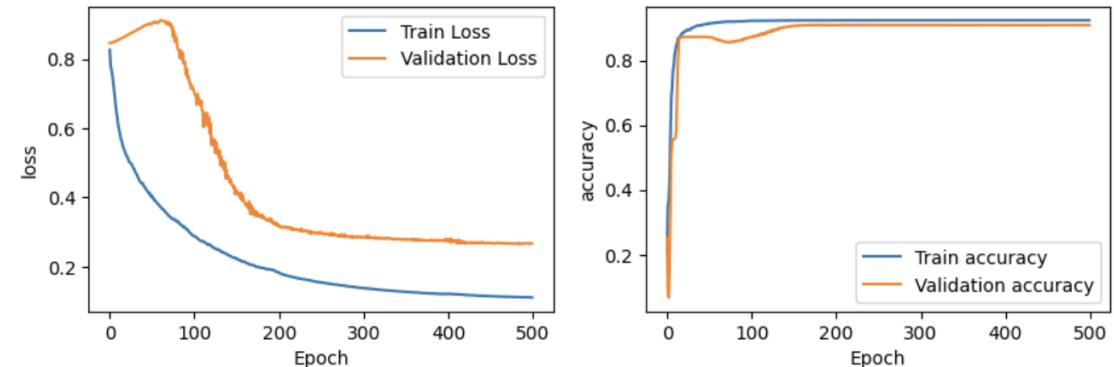
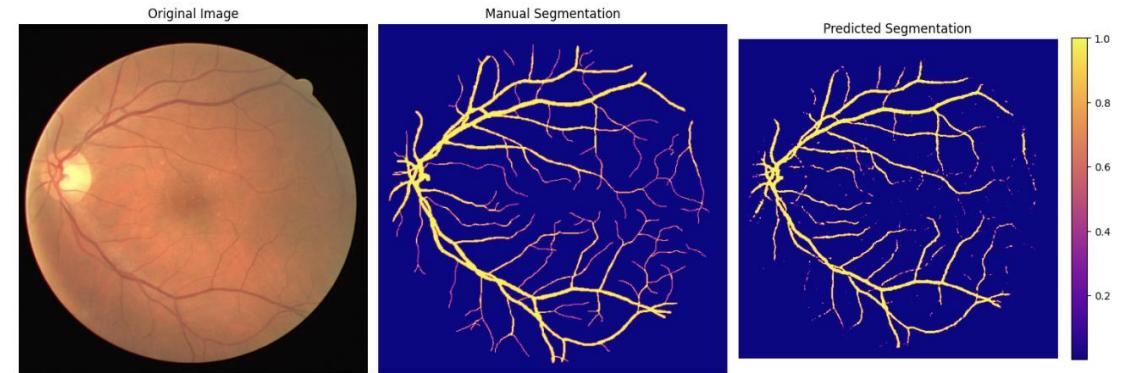
Deformable Encoder

- Validation DICE: 0.7433
- Parameters: 32,898,497
(trainable: 32,886,721)



Full DU-Net

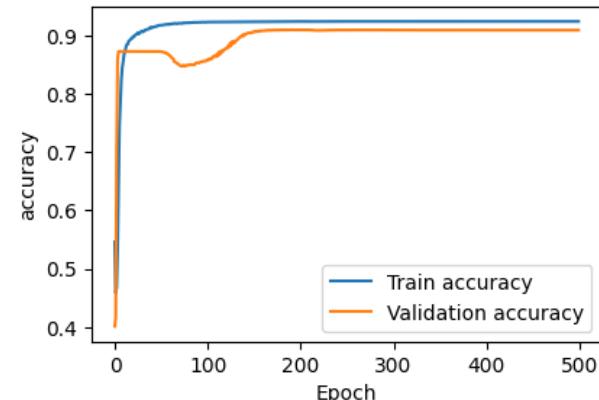
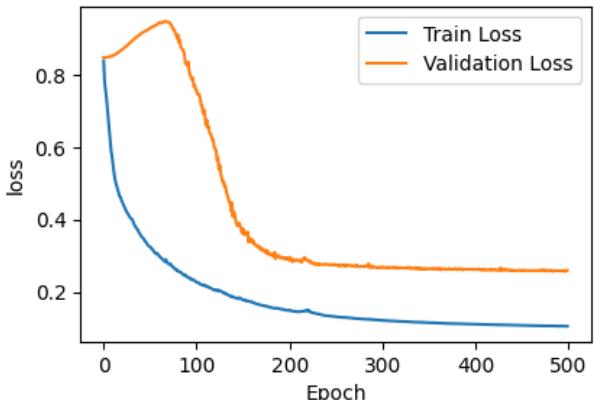
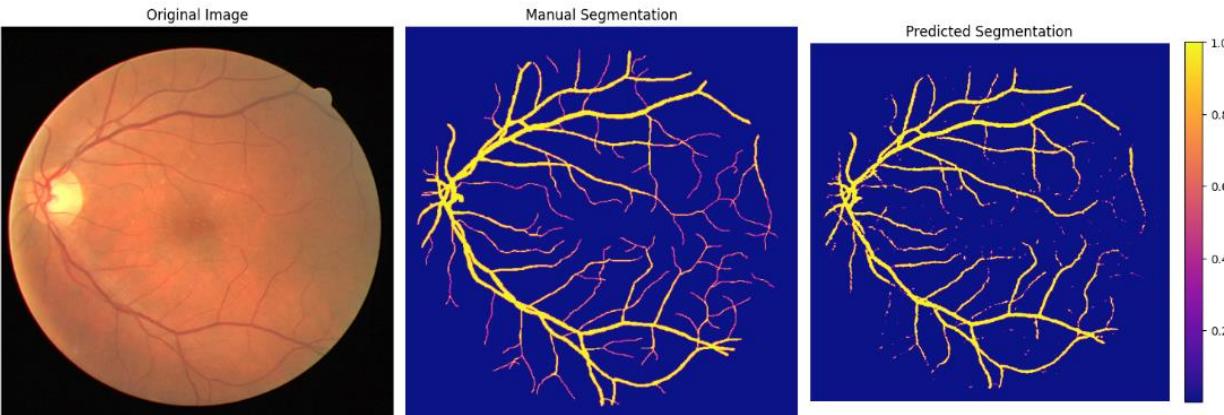
- Validation DICE: 0.7328
- Parameters: 34,741,697
(trainable: 34,729,921)



Performance

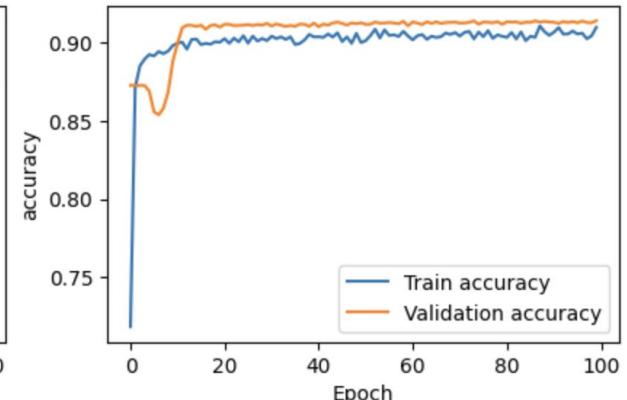
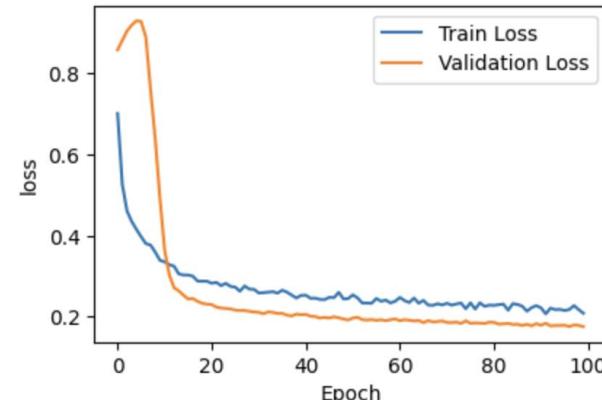
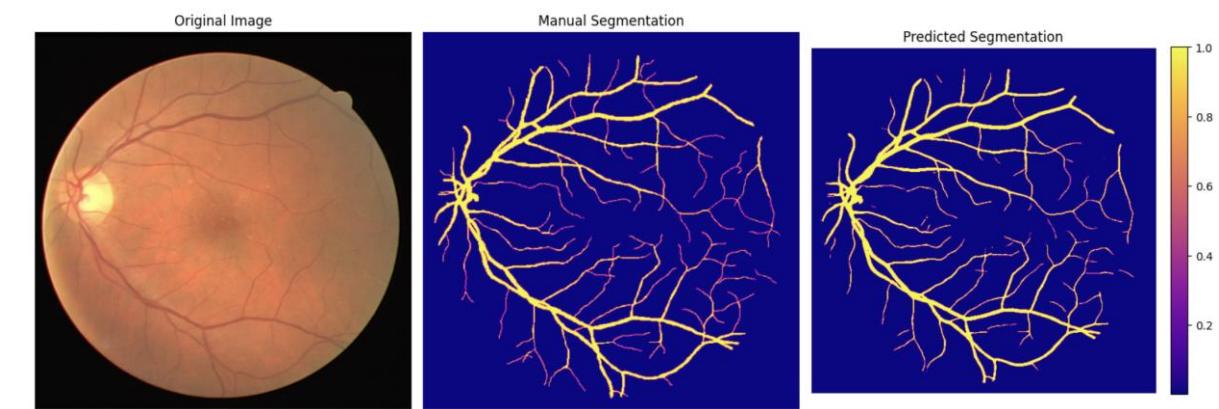
Deformable Encoder - Original Data

- Validation DICE: 0.7433



Deformable Encoder - Data Augmentation

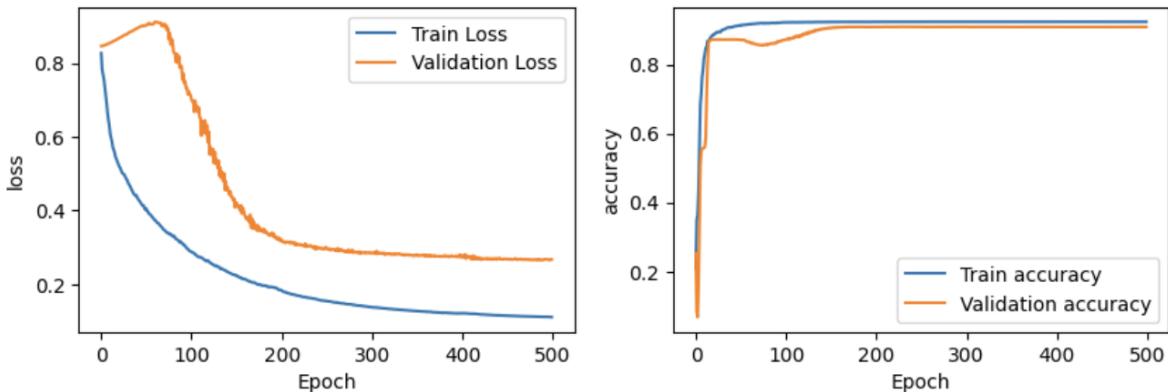
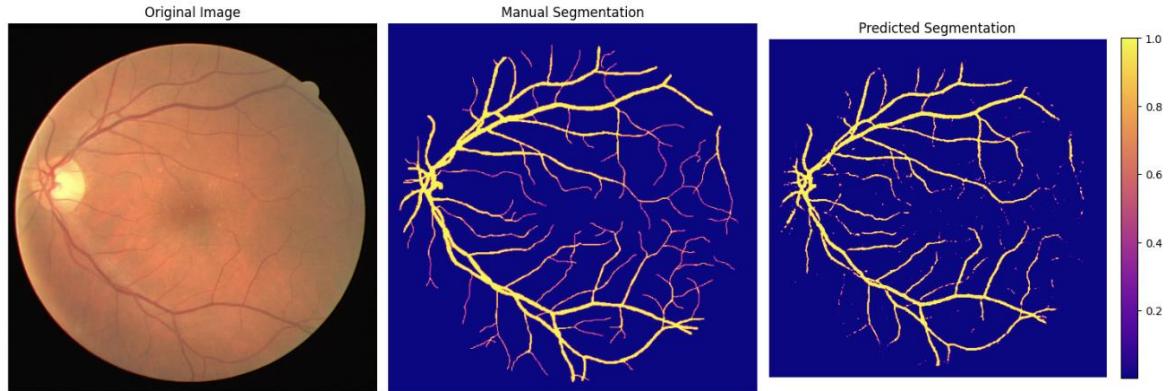
- Validation DICE: 0.8243



Performance

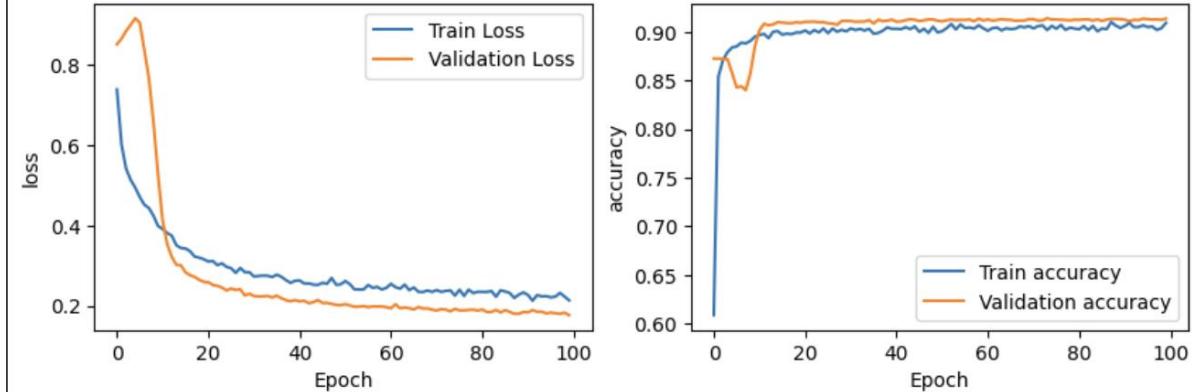
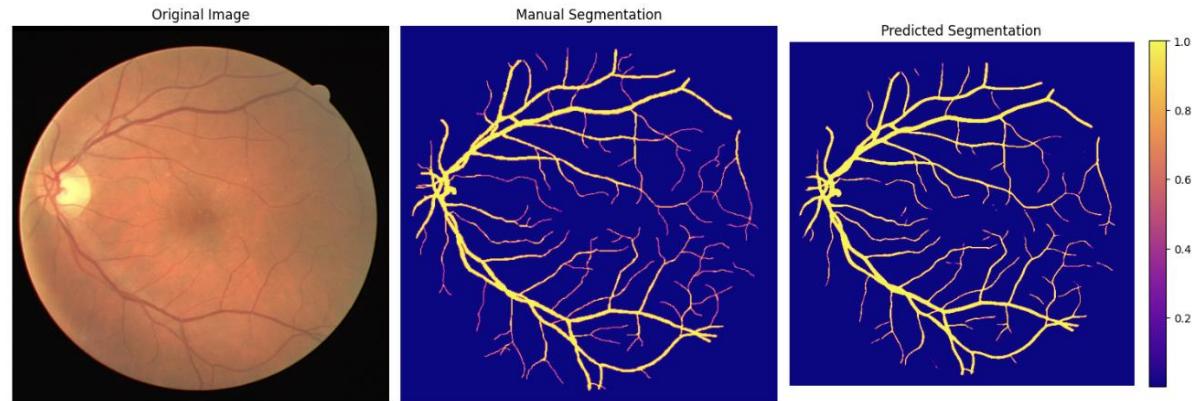
Full DUNet - Original Data

- Validation DICE: 0.7328

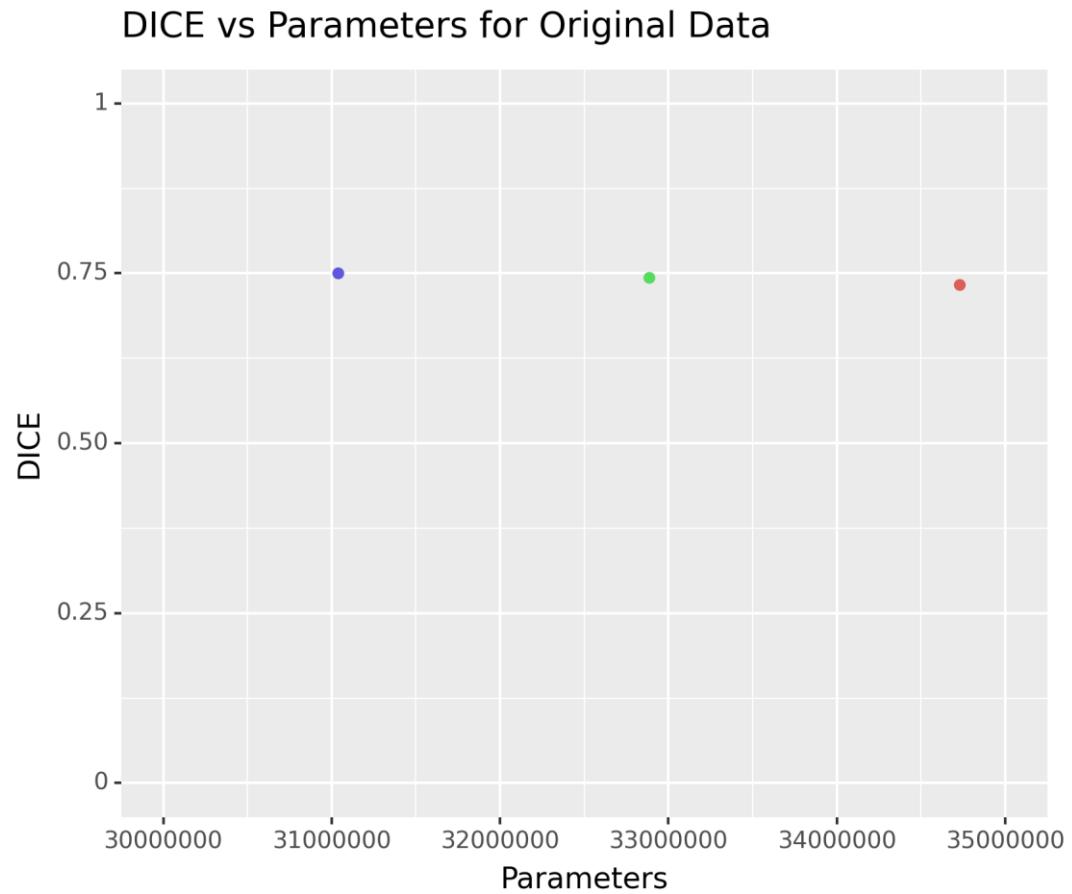


Full DUNet - Data Augmentation

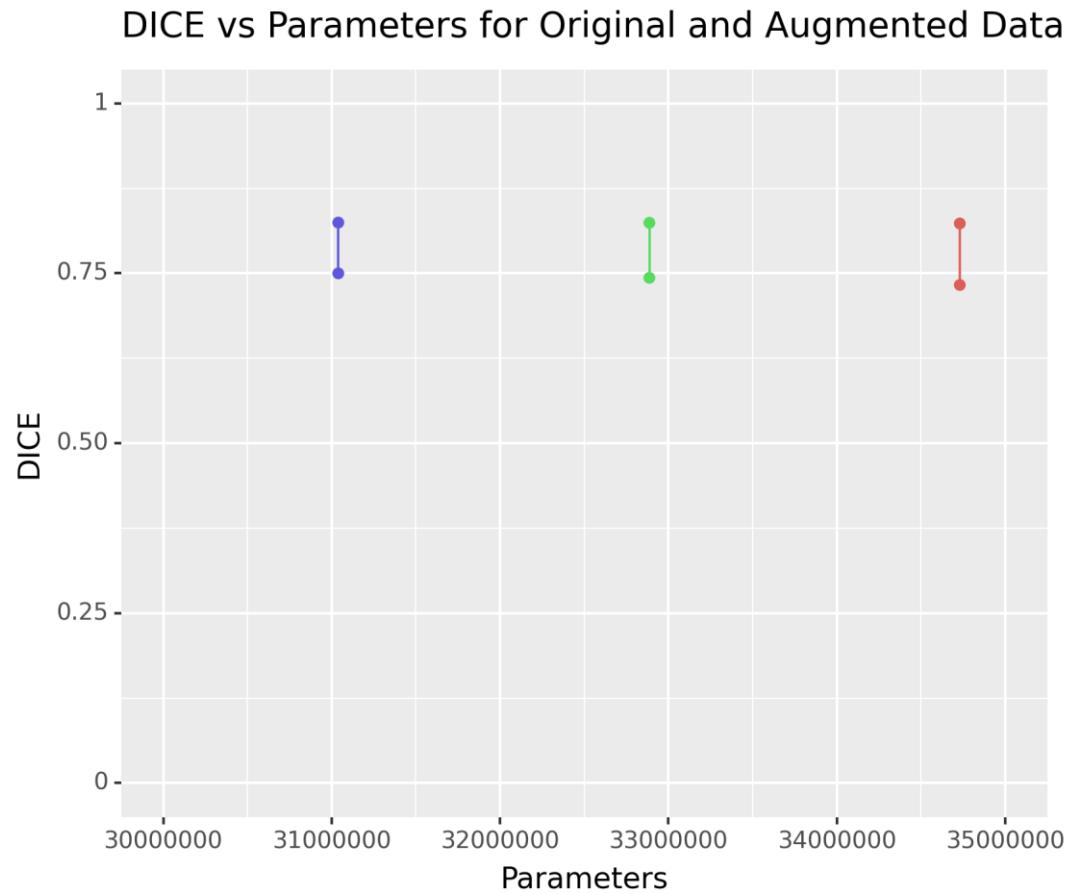
- Validation DICE: 0.8233



Effect of Data Augmentation

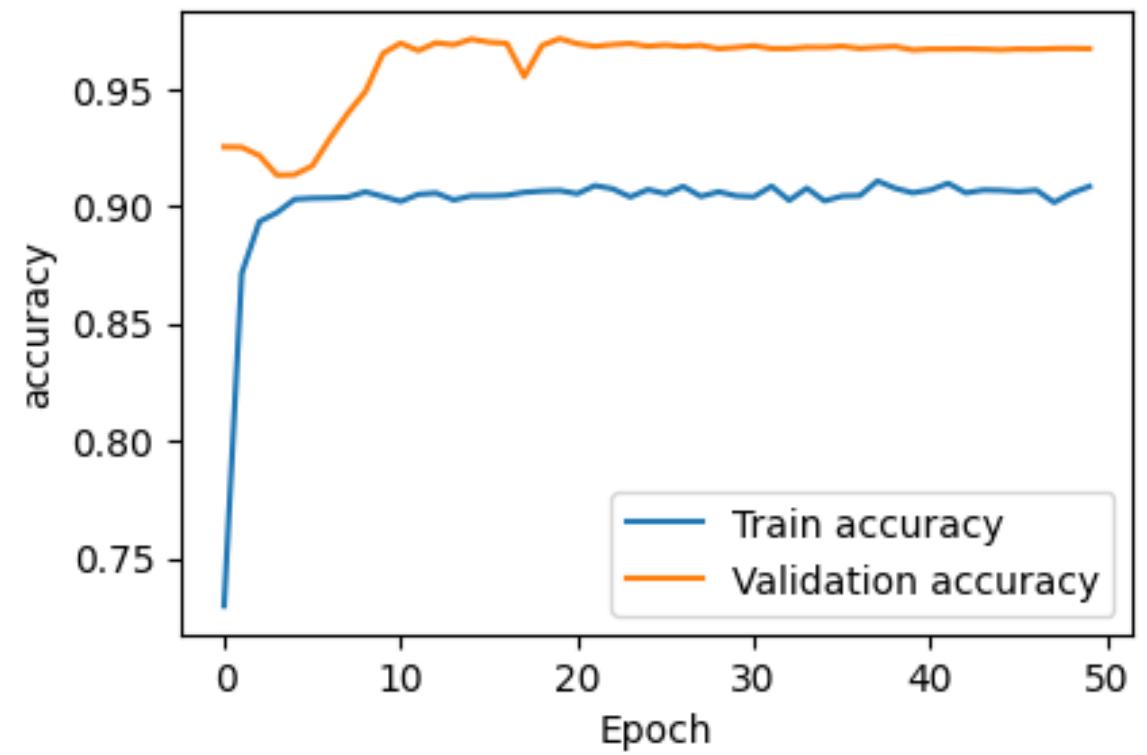
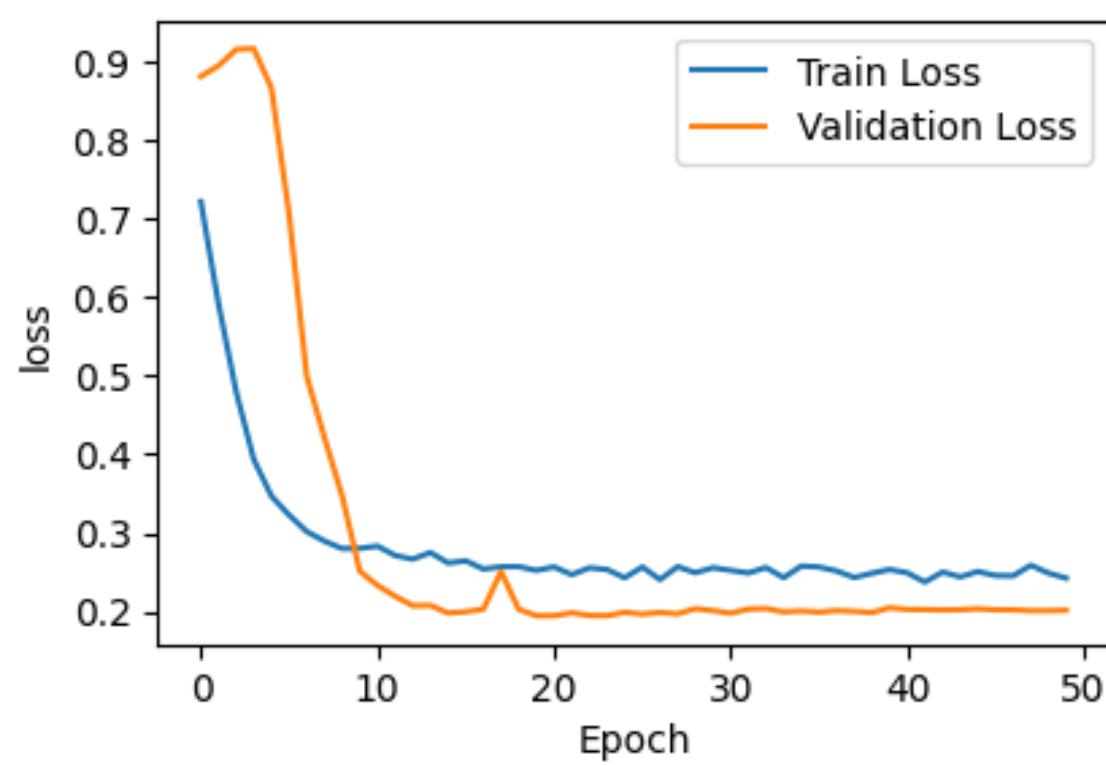


Effect of Data Augmentation



Results

Visualizations

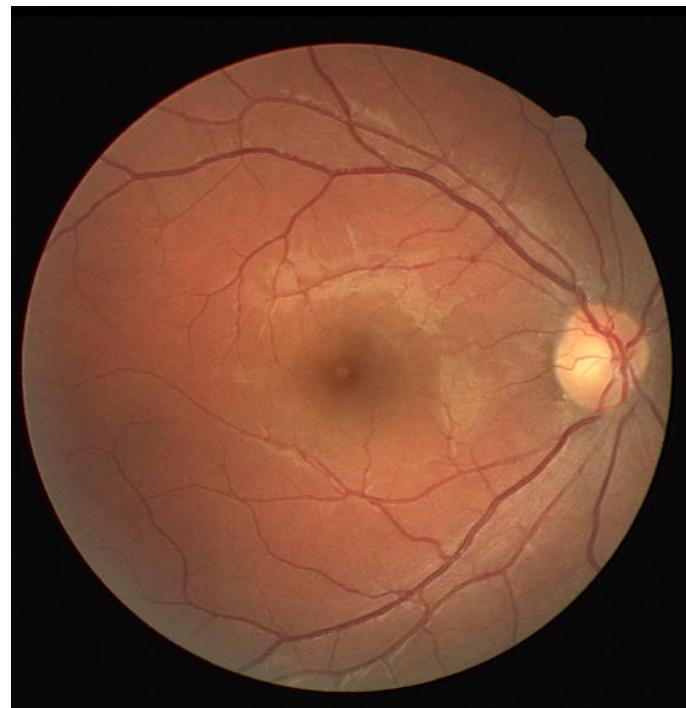
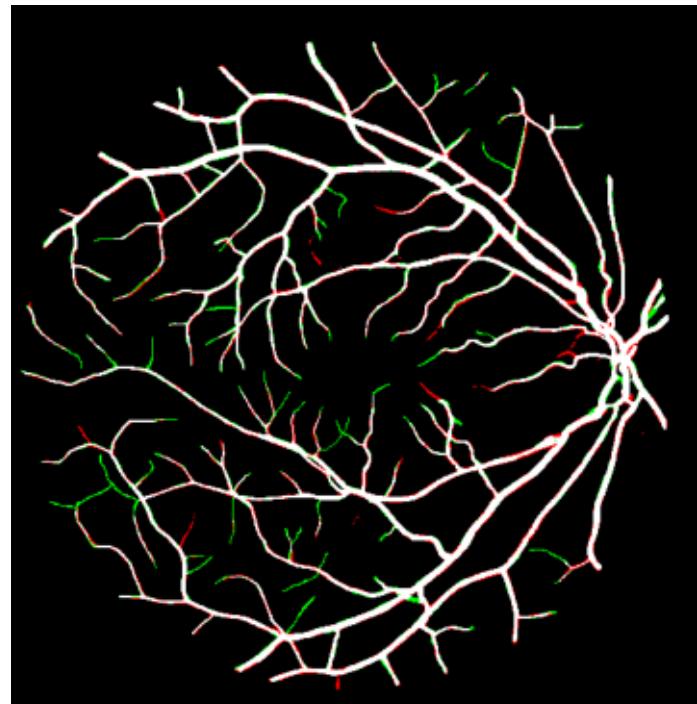


Results

CBAM-UNet (CFFA)

Test DICE: **0.8227**

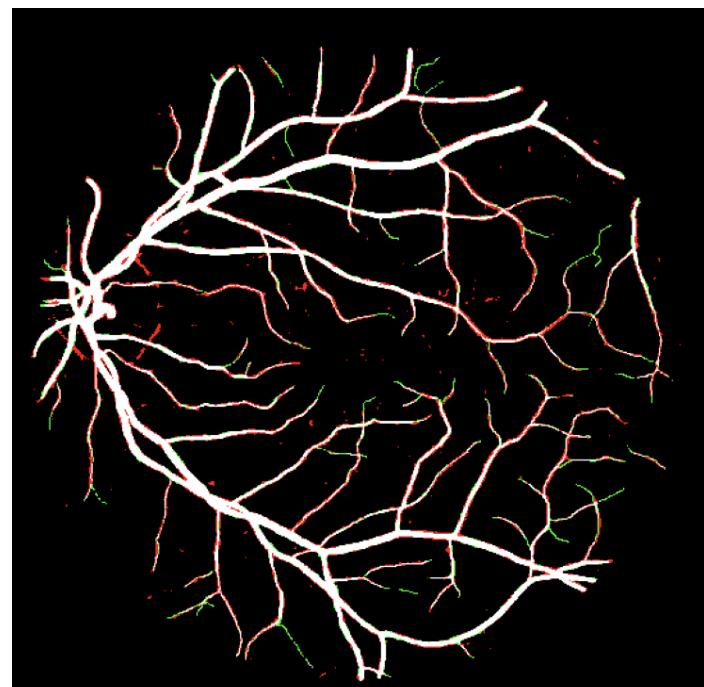
Grand Challenge rank: 306



TP	
TN	
FP	
FN	

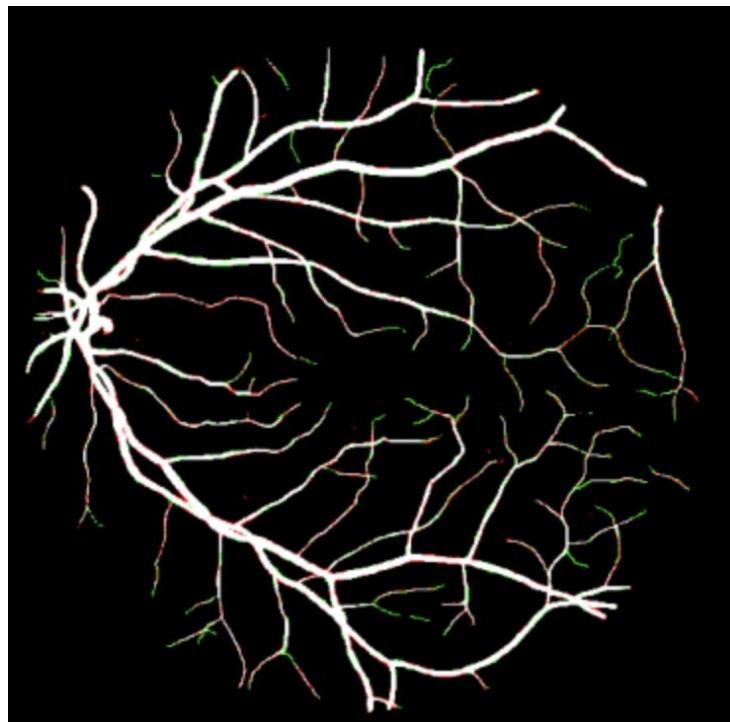
D-Unet

Validation DICE: **0.8233**



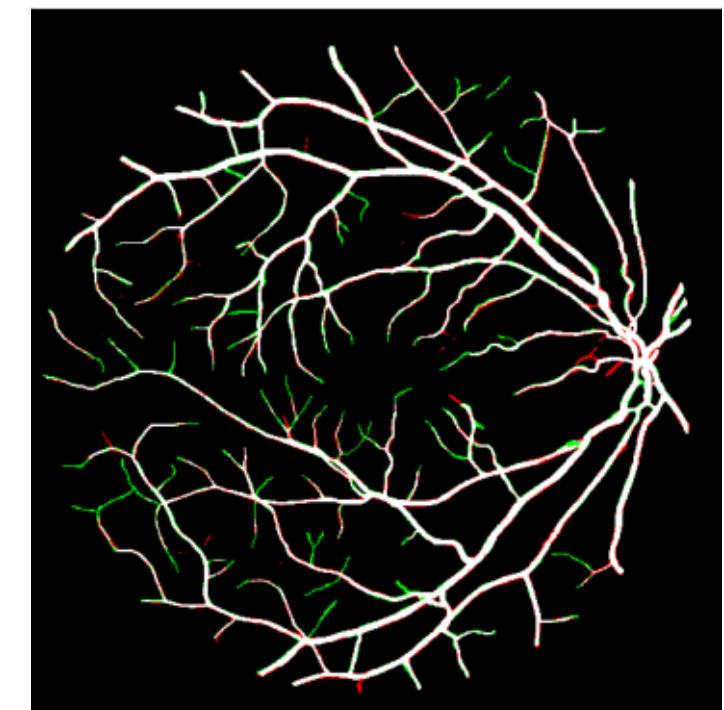
Results

UNet (3 encoders/decoders)
Test DICE: **0.8229**
Grand Challenge rank: 300



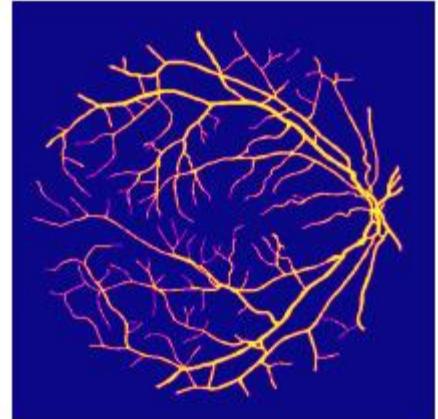
TP	
TN	
FP	
FN	

CMP-UNet
Test DICE: **0.8245**
Grand Challenge rank: 203

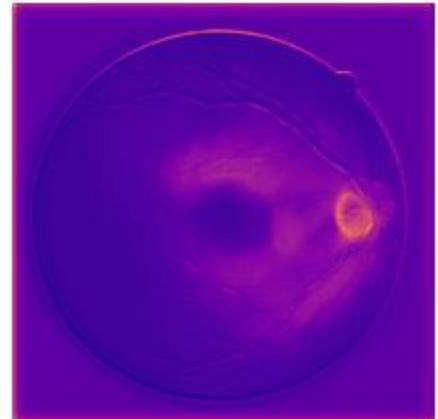


Epoch: 2

True

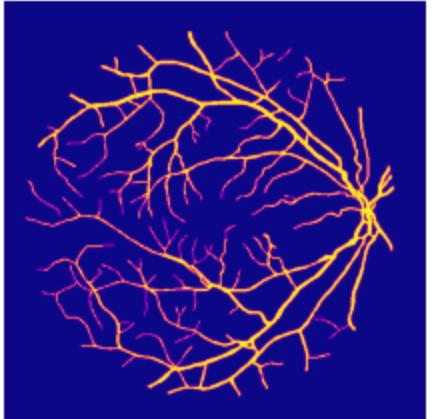


Pred

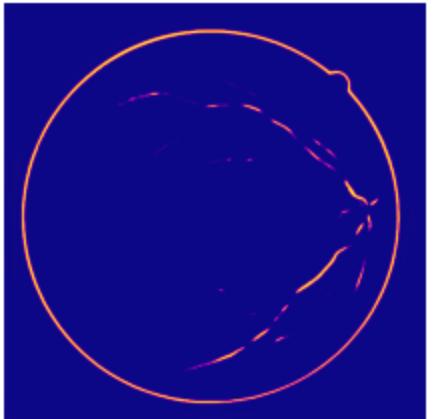


Epoch: 4

True

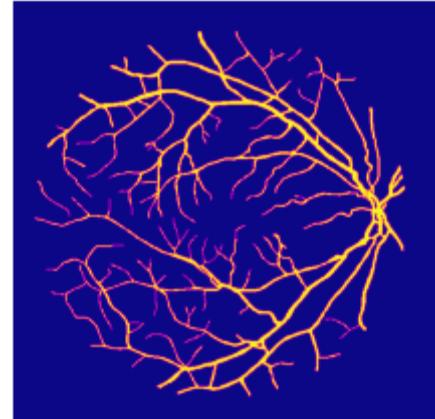


Pred

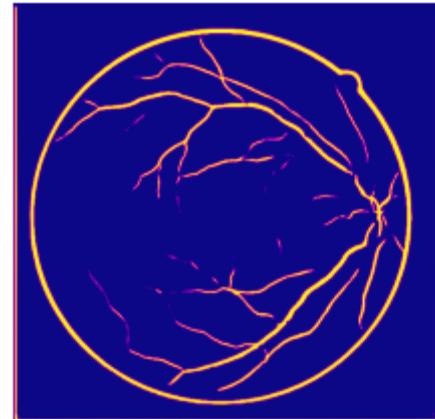


Epoch: 6

True

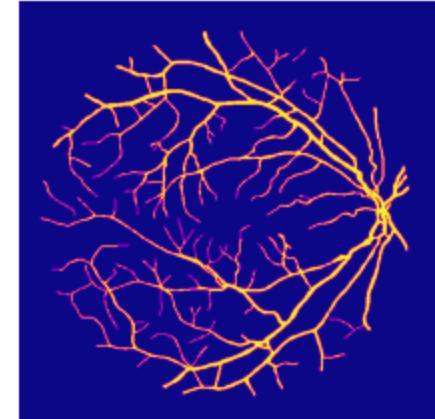


Pred

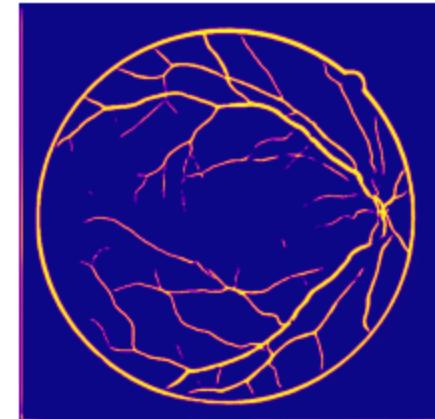


Epoch: 8

True

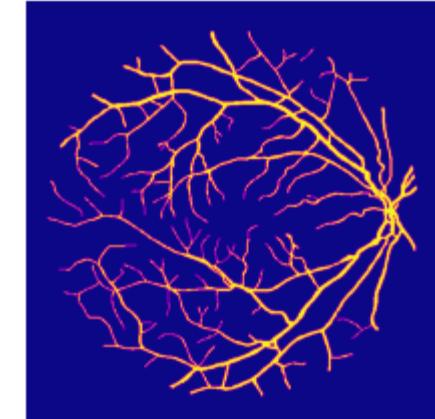


Pred

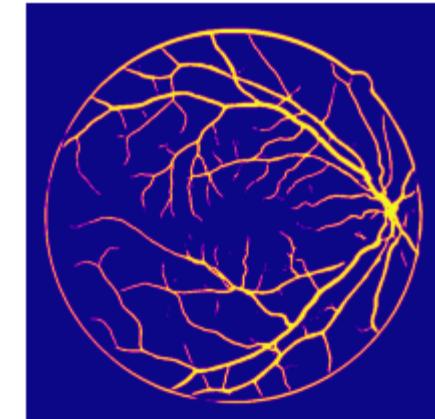


Epoch: 10

True



Pred



Bibliography

- Gu Y, Cao R, Wang D, Lu B. CMP-UNet: A Retinal Vessel Segmentation Network Based on Multi-Scale Feature Fusion. *Electronics*. 2023; 12(23):4743.
<https://doi.org/10.3390/electronics12234743>
- Guo, C., Szemenyei, M., Yi, Y., Wang, W., Chen, B., & Fan, C. (2020). SA-UNet: Spatial Attention U-Net for Retinal Vessel Segmentation. ArXiv, abs/2004.03696.
- Woo, S., Park, J., Lee, J.-Y., & Kweon, I. S. (2018). CBAM: Convolutional Block Attention Module.

Timesheet

Name	Subject	Estimated Duration
Simon De Lange	Base UNet and helper functions	4h
	CMP-UNet	16h
	CBAM-UNet (and extra modules)	6h
	Slides CBAM-UNet + CMP-UNet	2h
Mariska Van de Sompele	DUNet (custom layer)	16h
	DUNet implementations	3h
	Slides and Illustrations	3h
Jente Nijs	Background research	3h
	Data augmentation slide	2h